

Linux Memory In Practice

May 13, 2016

1 Typical memory issues

- linux ate my RAM
- kernel panic

2 Practise

- mystery under malloc
- is it a valid kernel addr
- detect memory overflow
- dynamic-size buffer
- persistent memory

3 More info

- learning resource

Low on Memory

Something like:

```
char *p = malloc(100);  
may fail...ENOMEM
```

.
.
.

(/proc/sys/vm/overcommit_memory)

Even worse, Out of Memory

```
md 265289728 sc: page allocation failure. order:7, mode:0xd0
```

Call Trace:

```
[<fffffffffc100e520>] dump_stack+0x8/0x34
```

```
[<fffffffffc10c95cc>] __alloc_pages_nodemask+0x5f4/0x700
```

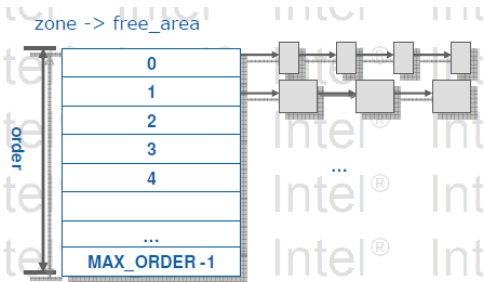
```
[<ffffffffffe037165c>] async_read_data_from_md+0x264/0x520 [vd]
```

```
[<ffffffffffe0362d9c>] md_demo+0x14c/0x4e0 [vd]
```

```
[<fffffffffc1080648>] kthread+0x88/0x90
```

```
[<fffffffffc1028a50>] kernel_thread_helper+0x10/0x18
```

Even worse, Out of Memory



memory usage from sysrq

Note: This is a similar output to OOM..

```
Node 0 Normal free:752948kB min:2464kB low:3080kB high:3696kB
active_anon:189164kB inactive_anon:118160kB active_file:64912kB
inactive_file:58544kB unevictable:0kB isolated(anon):0kB
isolated(file):0kB present:1308672kB managed:1241844kB mlocked:0kB
dirty:4kB writeback:0kB mapped:61264kB shmem:118564kB
slab_reclaimable:14688kB slab_unreclaimable:12808kB
kernel_stack:3024kB pagetables:8332kB unstable:0kB bounce:0kB
free_pcp:2492kB local_pcp:464kB free_cma:0kB writeback_tmp:0kB
pages_scanned:0 all_unreclaimable? no
```

Find the worst offenders

Who was blamed for it?

1. How much RAM has kernel used?
2. How much RAM has user space used?

RAM used by kernel - free

```
$ free
```

	total	used	free	shared	buff/cache
Mem:	3950132	626020	2081976	428552	1242136
Swap:	3906556	0	3906556		

$\text{total} = \text{used} + \text{free} + \text{buff/cache}$

RAM used by kernel - free

	total	used	free	shared	buffers	cached
Mem:	32065	7931	24133	6	353	6386
-/+ buffers/cache:		1191	30874			
Swap:	32688	0	32688			

approximately

used - buff/cache: $7931 - 6386 - 353$

free + buff/cache: $24133 + 6386 + 353$

RAM used by kernel - /proc/meminfo

```
$ cat /proc/meminfo
MemTotal:          3950132 kB
MemFree:           2083432 kB
MemAvailable:      2608652 kB
Buffers:           136580 kB
Cached:            976440 kB
SwapCached:         0 kB
Active:            877960 kB
Inactive:          788264 kB
Active(anon):      554928 kB
Inactive(anon):    426840 kB
Active(file):       323032 kB
Inactive(file):    361424 kB
Unevictable:        32 kB
Mlocked:            32 kB
SwapTotal:         3906556 kB
```

RAM used by kernel - /proc/meminfo

Slab:	128268 kB
SReclaimable:	93372 kB
SUnreclaim:	34896 kB
KernelStack:	6896 kB
PageTables:	27064 kB
NFS_Unstable:	0 kB
Bounce:	0 kB
WritebackTmp:	0 kB
CommitLimit:	5881620 kB
Committed AS:	3614800 kB
VmallocTotal:	34359738367 kB
VmallocUsed:	0 kB
VmallocChunk:	0 kB
HardwareCorrupted:	0 kB
AnonHugePages:	268288 kB
CmaTotal:	0 kB

RAM used by kernel - /proc/meminfo

```
VmallocUsed:          0 kB
VmallocChunk:          0 kB
HardwareCorrupted:     0 kB
AnonHugePages:        268288 kB
CmaTotal:              0 kB
CmaFree:               0 kB
HugePages_Total:       0
HugePages_Free:         0
HugePages_Rsvd:         0
HugePages_Surp:         0
Hugepagesize:          2048 kB
DirectMap4k:           108008 kB
DirectMap2M:           1892352 kB
DirectMap1G:           3145728 kB
```

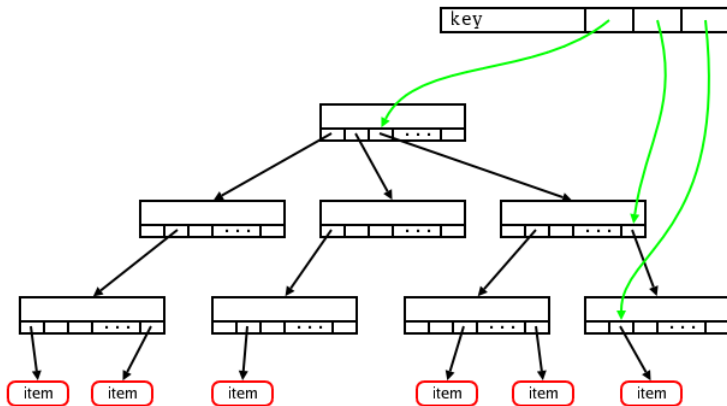
RAM used by kernel - slab?

slabtop

OBJS	ACTIVE	USE	OBJ	SIZE	SLABS	OBJ/SLAB	CACHE	SIZE	NAME
81186	81186	100%	0.19K	3866	21	15464K	dentry		
58188	58188	100%	0.10K	1492	39	5968K	buffer_head		
48174	48174	100%	1.01K	1554	31	49728K	ext4_inode_cache		
31110	31110	100%	0.12K	915	34	3660K	kernfs_node_cache		
27840	27582	99%	0.06K	435	64	1740K	kmalloc-64		
23400	23273	99%	0.20K	1170	20	4680K	vm_area_struct		
17442	17442	100%	0.04K	171	102	684K	ext4_extent_status		
15232	15232	100%	0.55K	544	28	8704K	inode_cache		
14208	13615	95%	0.03K	111	128	444K	kmalloc-32		
8568	8568	100%	0.08K	168	51	672K	anon_vma		

RAM used by kernel - page cache?

page cache: radix tree



RAM used by kernel - vmalloc?

```
cat /proc/vmallocinfo
```

```
0xfffffc9000000000-0xfffffc9000000200 8192 hpet_enable.part.13+0x1c/0x2a5 phys=fed000
0xfffffc9000000200-0xfffffc9000000400 8192 bpf_prog_alloc+0x35/0xa0 pages=1 vmalloc M
0xfffffc9000000400-0xfffffc9000000700 12288 acpi_os_map_iomem+0xf3/0x151 phys=aa405000
0xfffffc9000000800-0xfffffc9000000e00 24576 acpi_os_map_iomem+0xf3/0x151 phys=aa406000
0xfffffc9000000e00-0xfffffc9000001000 8192 acpi_os_map_iomem+0xf3/0x151 phys=abb4e000
0xfffffc9000001000-0xfffffc9000001f00 61440 acpi_os_map_iomem+0xf3/0x151 phys=aa3f8000
0xfffffc9000001f00-
0xfffffc90000042000 4198400 alloc_large_system_hash+0x160/0x221 pages=1024 vmalloc vpages
0xfffffc90000042000-0xfffffc900000423000 12288 alloc_large_system_hash+0x160/0x221 pages=
0xfffffc900000423000-0xfffffc900000624000 2101248 alloc_large_system_hash+0x160/0x221 pages=
```

```
cat /proc/vmallocinfo
```

RAM used by process - top

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3769	root	20	0	49012	4156	3408	R	0.3	0.1	0:02.14	top
1	root	20	0	119912	6004	3924	S	0.0	0.2	0:02.09	system
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthrea
3	root	20	0	0	0	0	S	0.0	0.0	0:00.01	ksofti
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworke


top + shift + m: Sort by memory usage

RAM used by process - /proc/pid/statm

```
cat /proc/3769/statm  
2253 1039 852 24 0 351 0
```

$$1039 * 4 = 4156k$$

RAM used by process - process address space



RAM used by process - memory layout

```
0000000000000000 - 00007fffffffffff (=47 bits) user space
ffff880000000000 - ffffc7fffffffffff (=64 TB) direct mapping of all mem
ffffc90000000000 - ffffe8fffffffffff (=45 bits) vmalloc/ioremap space
ffffea0000000000 - ffffeafffffffffffff (=40 bits) virtual memory map (1TB)
ffffffff80000000 - fffffffffa0000000 (=512 MB) kernel text mapping
fffffffffa0000000 - ffffffff5fffff (=1526 MB) module mapping space
```

RAM used by process - /proc/pid/maps

Typical process addresss space:

```
cat /proc/3769/maps
00400000-00418000 r-xp 00000000 08:05 1050669
00619000-0061b000 rw-p 00019000 08:05 1050669
0157d000-01624000 rw-p 00000000 00:00 0
7ffed889f000-7ffed88c0000 rw-p 00000000 00:00 0
```

heap

RAM used by process - /proc/pid/smmaps

```
0157d000-01624000 rw-p 00000000 00:00 0
```

```
Size: 668 kB
```

```
Rss: 404 kB
```

```
Pss: 404 kB
```

```
Shared_Clean: 0 kB
```

```
Shared_Dirty: 0 kB
```

```
Private_Clean: 0 kB
```

```
Private_Dirty: 404 kB
```

```
Referenced: 404 kB
```

```
Anonymous: 404 kB
```

```
AnonHugePages: 0 kB
```

```
Shared_Hugetlb: 0 kB
```

```
Private_Hugetlb: 0 kB
```

```
Swap: 0 kB
```

```
SwapPss: 0 kB
```

```
KernelPageSize: 4 kB
```

```
MMUPageSize: 4 kB
```

```
Locked: 0 kB
```

```
VmFlags: rd wr mr mw me ac sd
```

kernel detect critical issue

slab caused kernel panic

slab bug_on triggered

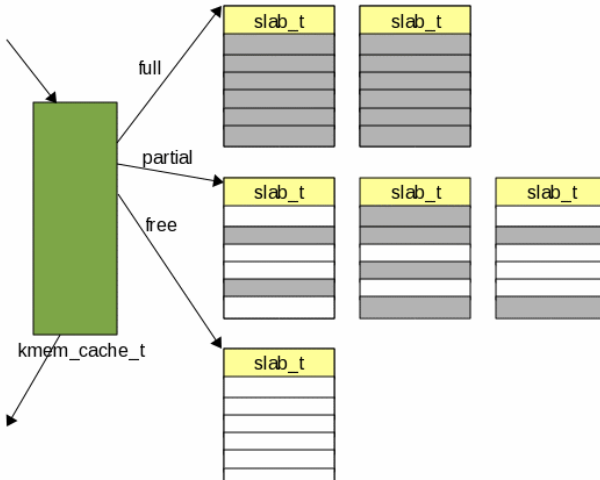
```
kernel BUG at mm/slab.c:3067!  
Unable to handle kernel NULL pointer dereference at virtual address 00000000  
[<c003c6ec>] (__dabt_svc+0x4c/0x60) from [<c00405a4>] (__bug+0x1c/0x28)  
[<c00405a4>] (__bug+0x1c/0x28) from [<c00c9d0c>] (cache_alloc_refill+0x3a0/0x654)  
[<c00c9d0c>] (cache_alloc_refill+0x3a0/0x654) from [<c00ca174>] (kmem_cache_alloc+0xb0/0xc4)  
[<c00ca174>] (kmem_cache_alloc+0xb0/0xc4) from [<c0057b04>] [color=Red](copy_process+0x9c/0xdbc)  
[<c0057b04>] (copy_process+0x9c/0xdbc) from [<c0058890>] (do_fork+0x48/0x288)  
[<c0058890>] (do_fork+0x48/0x288) from [<c003cc40>] (ret_fast_syscall+0x0/0x30)
```

do_fork : copy_process : kmem_cache_alloc : cache_alloc_refill :
__bug


slab bug_on triggered

```
/*  
 * The slab was either on partial or free list so  
 * there must be at least one object available for  
 * allocation.  
 */  
BUG_ON(slabp->inuse >= cachep->num);
```


slab structure




slab objects



slab objects inside

Slab with Internal Descriptors



dump of slab

Hexdump:

```
000: 00 50 90 df e0 2d 80 df 40 00 00 00 40 30 91 df
010: 05 00 00 00 ff fe ff ff 00 00 ad de 03 00 00 00
020: ff ff ff ff 00 00 00 00 ff fe ff ff ff ff ff ff
```

slab info

slab.free = ffffffff means there is no free obj in this slab, but
fffffeff?

```
struct slab {  
    union {  
        struct {  
            struct list_head list; 0xdf905000,0xdf80  
            unsigned long colouroff; 0x40  
            void *s_mem; /* including  
            unsigned int inuse; /* num of objs  
            kmem_bufctl_t free; 0xfffffeff typedef un  
            unsigned short nodeid;  
        };  
        struct slab_rcu __slab_cover_slab_rcu;  
    };  
};
```

malloc

What happened when:

```
char *p = malloc(100); // only virtual memory allocated
```

```
memset(p, 0, 100); // physical memory allocated
```

Is it a valid kernel address?

How?

`virt_addr_valid?`

Is it a valid kernel address?

```
0000000000000000 - 00007fffffffffff (=47 bits) user space
ffff880000000000 - fffffc7fffffffffff (=64 TB) direct mapping of all mem
ffffc90000000000 - fffffe8fffffffffff (=45 bits) vmalloc/ioremap space
fffffea000000000 - fffffeafffffffffffff (=40 bits) virtual memory map (1TB)
fffffffff80000000 - fffffffffa0000000 (=512 MB) kernel text mapping
ffffffffffa0000000 - ffffffffffff5fffff (=1526 MB) module mapping space
```

kernel layout


Is it a valid kernel address?

Is it also a safe address to access?

A pagetable walk is required

Is it a valid kernel address?

page walker:



detect memory overflow for page/slab

How to detect memory overflow/overwrite for page/slab?

current implementation in kernel:

1. Post-detection

2. Runtime-detection

Post-detection


How?

SLAB_POISON : avoid double-free

SLAB_RED_ZONE : overflow detection

Post-detection

Object Format:



Runtime-detection

How?

Documentation/kmemcheck.txt

```
pages = alloc_pages(GFP_KERNEL,1); addr =  
page_address(pages); if(*addr == 'a' ) //kmemcheck warn.
```

kmemcheck is based on page fault and uses extra shadow page

Only support x86

Runtime-detection

Can we write a platform-independent detection method?

Runtime-detection

Let's learn from kmemcheck, we use a rb-tree to maintain all the allocated regions, and add checking-hooks in memcpy/memset/strcpy, etc

dynamical-size buffer

alloc_page? kmalloc? vmalloc?

ring buffer is fixed-size, and hard to use...

What can we learn from page cache?

convenient read(fd, offset, size)

What can we learn from page table?


only struct page is needed.

dynamical-size buffer

How about mixing the two together?

Let's turn page cache tree node into struct page.

dynamical-size buffer



$$4 * (4k / \text{sizeof}(\text{unsigned long}) * (4k / \text{sizeof}(\text{unsigned long}) * 4k \\ = 4.2G$$

persistent memory

How can we add a new zone to linux?
Check the git log of ZONE_DEVICE

Reference



memory layout.

https://www.kernel.org/doc/Documentation/x86/x86_64/mm.txt



Understanding the Linux Kernel, 3rd Edition.

<http://gauss.ececs.uc.edu/Courses/c4022/code/memory/understandi>



lwn

<https://lwn.net/Articles>



feel free to contact

yu.chen.surf@gmail.com

questions

Q/A