


CS353 Linux Kernel

Chentao Wu 吴晨涛
Associate Professor
Dept. of CSE, SJTU
wuct@cs.sjtu.edu.cn



上海交通大学

Introduction of Myself

■ Dual Ph.D.


- 2012, Electrical and Computer Engineering, *Virginia Commonwealth University (VCU)*, Richmond, VA, USA
- 2010, Computer Architecture, *Huazhong University of Science and Technology (HUST)*, Wuhan, China

■ Research Interest: **Data Storage Systems**

- Storage management for Big Data
- Cloud storage, Green storage
- Reliable storage systems (e.g., disk arrays)
- Semantic file systems (e.g., object-based storage sys.)
- Cache Algorithms in storage systems



上海交通大学




Our Lab

- Leader: Prof. Minyi Guo (Dean of CSE Dept.)
 - <http://epcc.sjtu.edu.cn>
 - **Parallel and Distributed Computing**
 - Parallel and Distributed Systems/Networks
 - High Performance Computing
 - Cloud Computing
 - Big Data
- Welcome to participate in our lab
- 15+ master students per year, 5+ doctoral students per year



上海交通大學



Download Lectures and Upload Projects

- <ftp://public.sjtu.edu.cn>
- User: wuct
- Password: wuct123456



上海交通大學



Teaching Assistant

- Chao Tan 谭超
 - Email: 345243921@qq.com
 - Mobile Phone: 15821274485
- Help me to review Projects and the Final Exam




上海交通大學



Books

- No textbook
- References
 - Understanding the Linux Kernel
 - 3rd Edition
 - Linux Kernel Drivers
 - 3rd Edition
 - Linux Kernel Development
 - 3rd Edition



上海交通大学




Syllabus (1)

■ Requirements:

- Computer Organization, Operating System
- C/C++ Programming

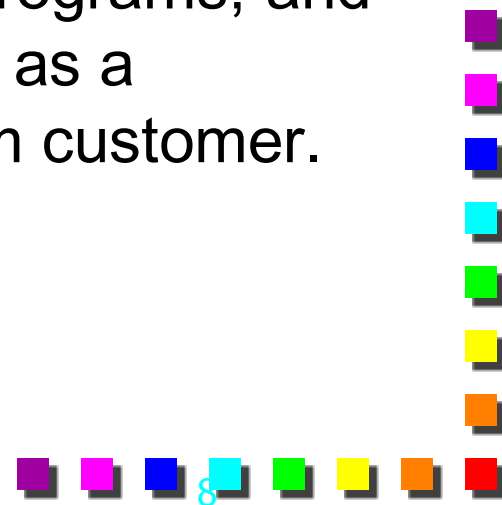
■ Goals: Successful course participants will:

- Understand C programming in Linux Kernel (Module programming in Linux Kernel).
- Understand the core concepts of operating systems, including processes, threads, synchronization, virtual memory policies, and file management.



Syllabus (2)

- Goals (contd.)
 - The idea of the course is **to learn how computers really work in Linux Kernel**, from the chip level up to the application level. When we finish, you will understand what is actually happening when a computer system is running a set of programs, and will be able to make informed choices as a developer, project manager, or system customer.




Course Meeting Times

■ Lectures:

- 2 classes per week (Friday)
- 2 classes per dual weeks (Virtual Time)

■ Questions:

- Ask me directly between/after the classes
- Go to my office: SEIEE 3-513
- Send me an email: wuct@cs.sjtu.edu.cn
- Ask teaching assistant




Final Grades

- Participation 10%
 - Randomly
- Project 20%
 - Four Projects
 - Source Code and Document
- Report (Reading Kernel Code or Update Experimental Manual) 20%
- Final Exam 50% (Close Book)




Late Policy

- Late Policy: Deadlines will be given in each assignment. These deadlines are strict.
- Typically, project will be given on Fridays, you should submit your homework by the next three weeks.



Cooperation

- This course is established with Intel-Shanghai Corporation.
 - Several Intel engineers will share their experiences on Linux Kernel Programming.



上海交通大學

1. The Linux Kernel: Introduction

Chentao Wu
Associate Professor
Dept. of CSE, SJTU
wuct@cs.sjtu.edu.cn



上海交通大學

What's Linux?

- **Linux** is a *Unix-like* and *POSIX-compliant* computer operating system assembled under the model of free and open source software development and distribution. The defining component of Linux is the **Linux kernel**, an operating system kernel first released on 5 October 1991 by **Linus Torvalds**.
(from <http://www.wikipedia.org>)




上海交通大學


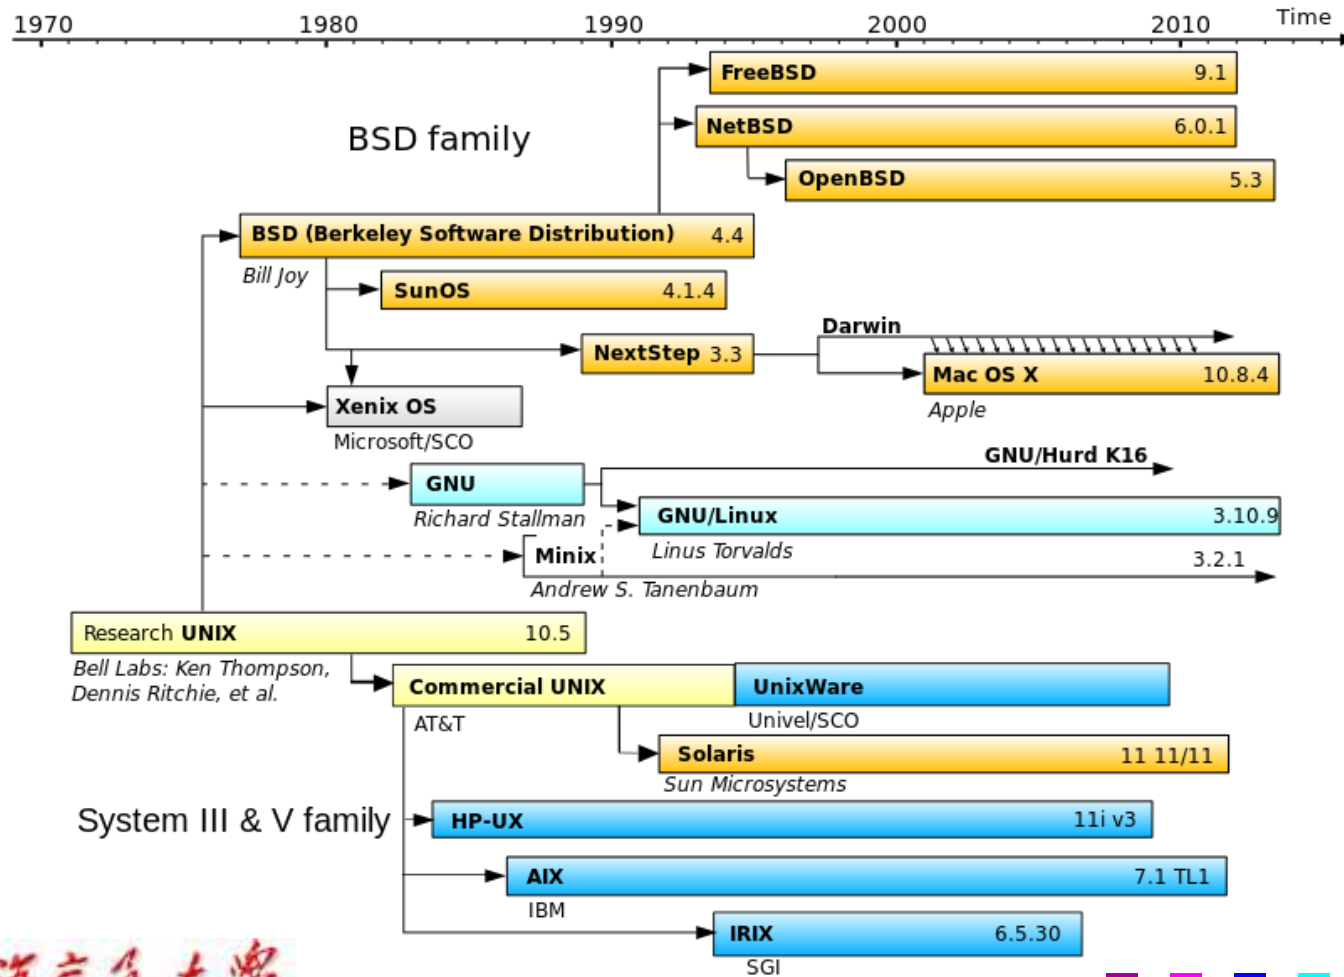


History

- UNIX: 1969 Thompson & Ritchie AT&T Bell Labs.
- BSD: 1978 Berkeley Software Distribution.
- Commercial Vendors: Sun, HP, IBM, SGI, DEC.
- GNU: 1984 Richard Stallman, FSF.
- POSIX: 1986 IEEE Portable Operating System unIX.
- Minix: 1987 Andy Tannenbaum.
- SVR4: 1989 AT&T and Sun.
- Linux: 1991 Linus Torvalds Intel 386 (i386).
- Open Source: GPL.



UNIX Family



上海交通大学

Linux Features

- UNIX-like operating system.
- Features:
 - Preemptive multitasking.
 - Virtual memory (protected memory, paging).
 - Shared libraries.
 - Demand loading, dynamic kernel modules.
 - Shared copy-on-write executables.
 - TCP/IP networking.
 - SMP support.
 - Open source.



Linux Distribution



上海交通大学



Ubuntu



上海交通大学


Two Modes within Linux

User mode	Application software (bash , LibreOffice , Blender , 0 A.D.)		
	Complex libraries (GLib , GTK+ , Qt , SDL , EFL)		Application software
	Complex libraries (GLib , GTK+ , Qt)	Simple libraries (sin , opendbm)	Application software
	open , exec , sbrk , socket , fopen , calloc C standard library: glibc (1,187,911 lines of code) / uClibc (342,842 lines of code)		
Kernel mode	System calls: TRAP , CALL , BRK , INT , IOCTL (depends on the hardware)		
	Linux kernel (16,223,920 lines of code) (device drivers, process-scheduler, networking stack, file systems) ALSA , DRI , evdev , LVM , device mapper , Linux Process Scheduler , Linux Network Scheduler , Netfilter Linux Security Modules: SELinux , TOMOYO , AppArmor , Smack		
	Hardware (CPU(s) , Memory , other Microprocessors , Devices etc)		



What's a Kernel?

- AKA: executive, system monitor.
- Controls and mediates access to hardware.
- Implements and supports fundamental abstractions:
 - Processes, files, devices etc.
- Schedules / allocates system resources:
 - Memory, CPU, disk, descriptors, etc.
- Enforces security and protection.
- Responds to user requests for service (system calls).
- Etc...etc...

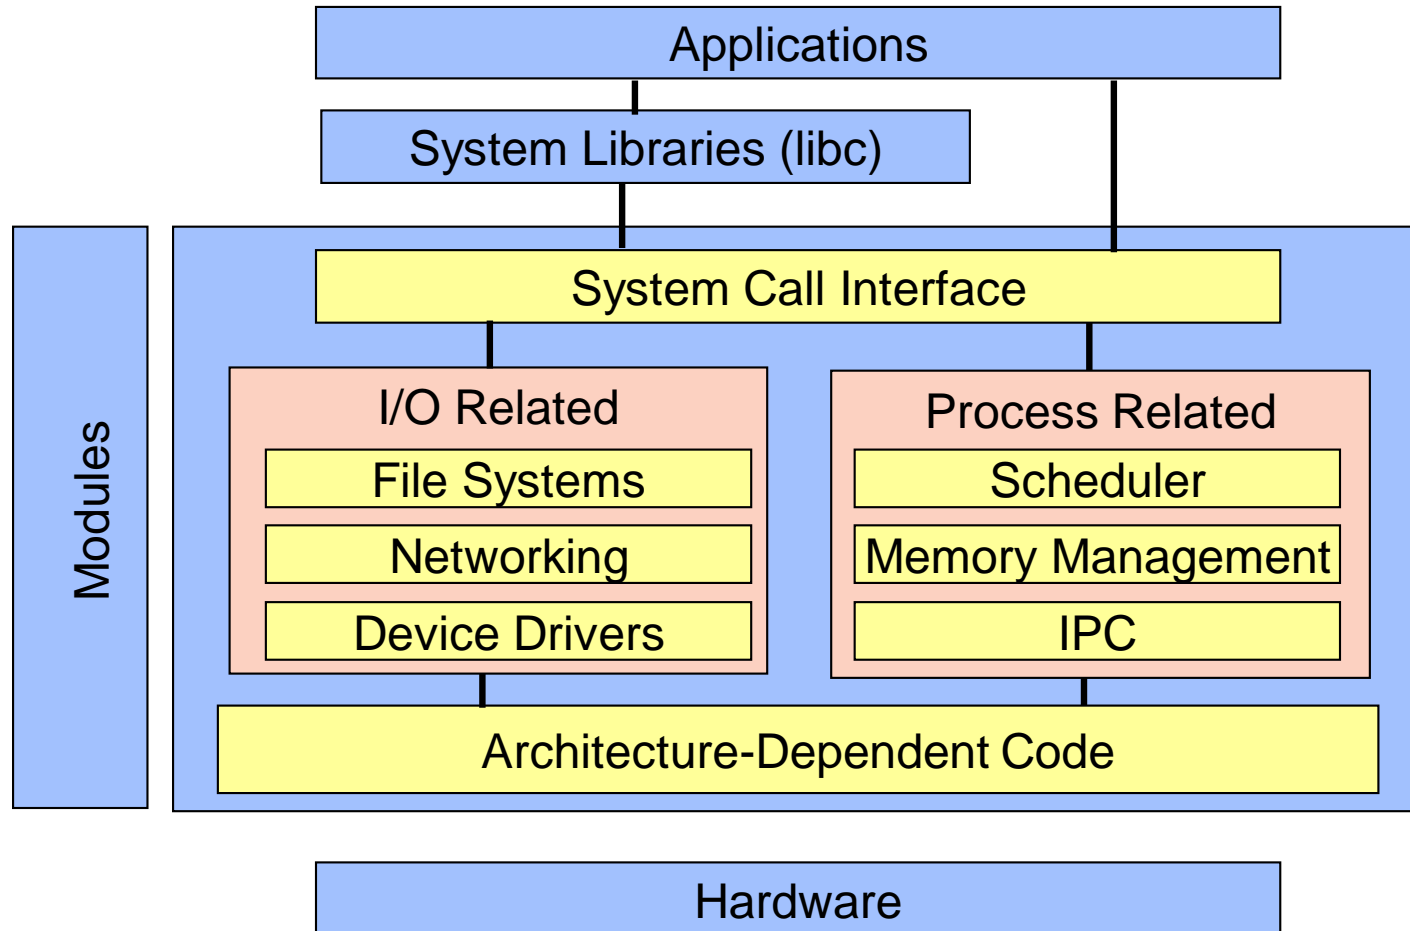


Kernel Design Goals

- Performance: efficiency, speed.
 - Utilize resources to capacity with low overhead.
- Stability: robustness, resilience.
 - Uptime, graceful degradation.
- Capability: features, flexibility, compatibility.
- Security, protection.
 - Protect users from each other & system from bad users.
- Portability.
- Extensibility.



Example “Core” Kernel

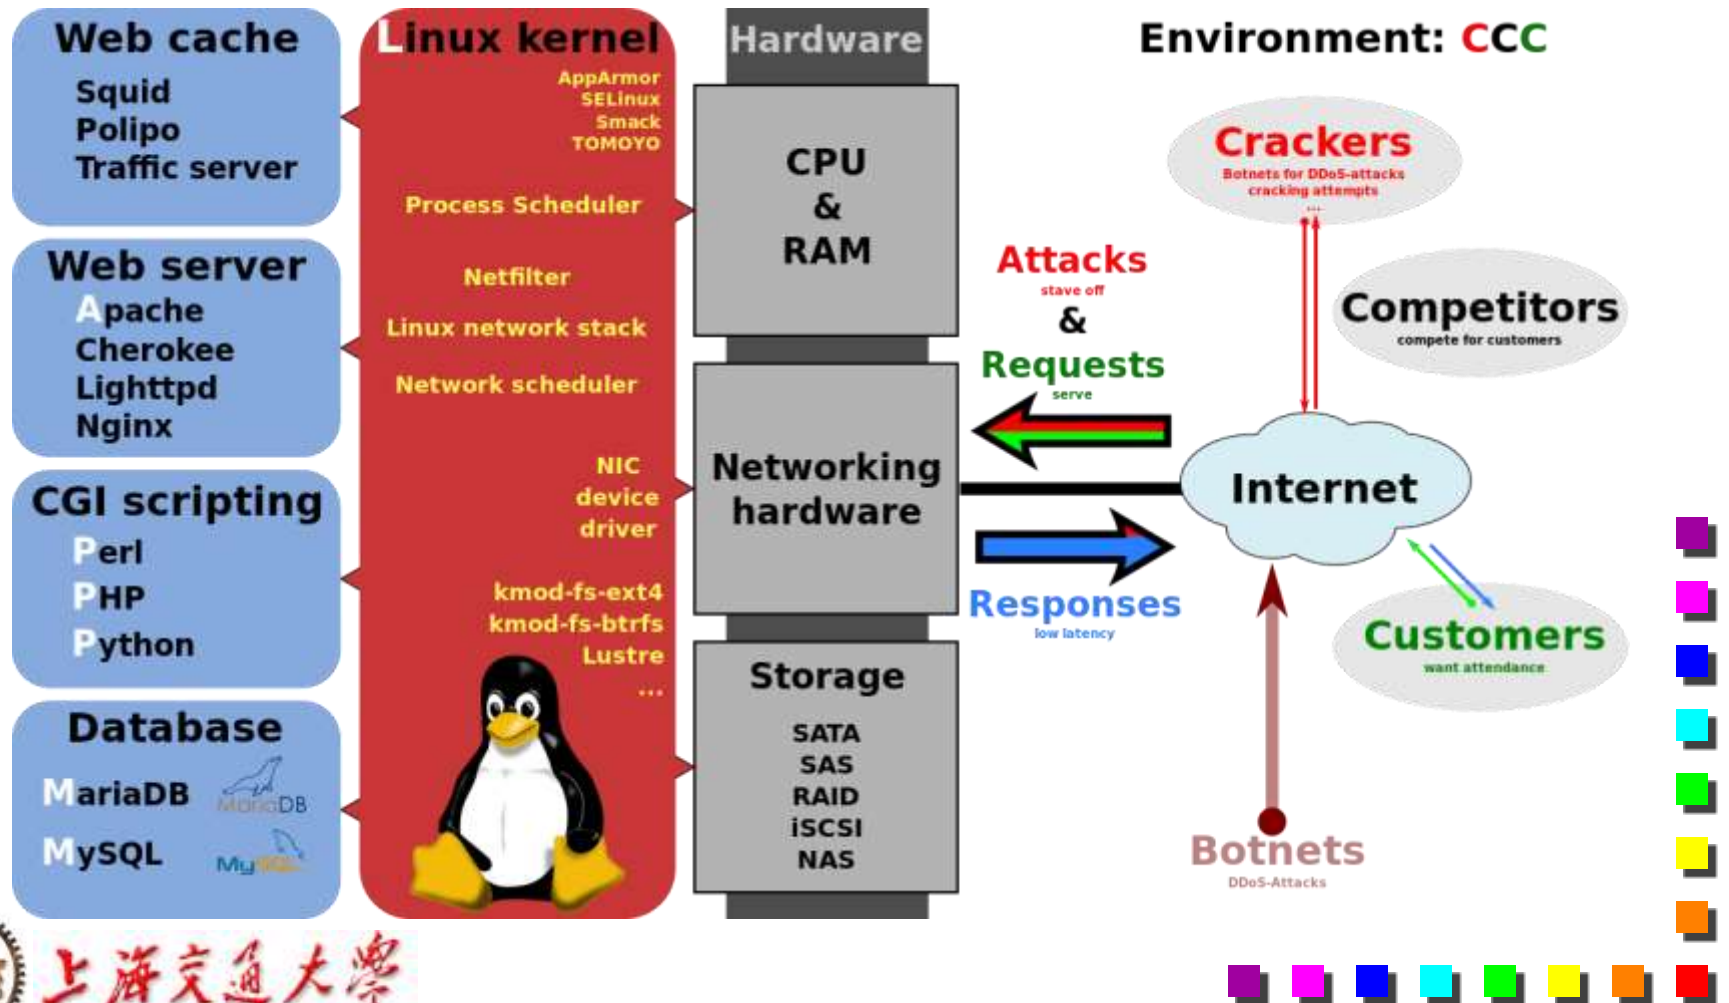


Architectural Approaches

- Monolithic.
- Layered.
- Modularized.
- Micro-kernel.
- Virtual machine.

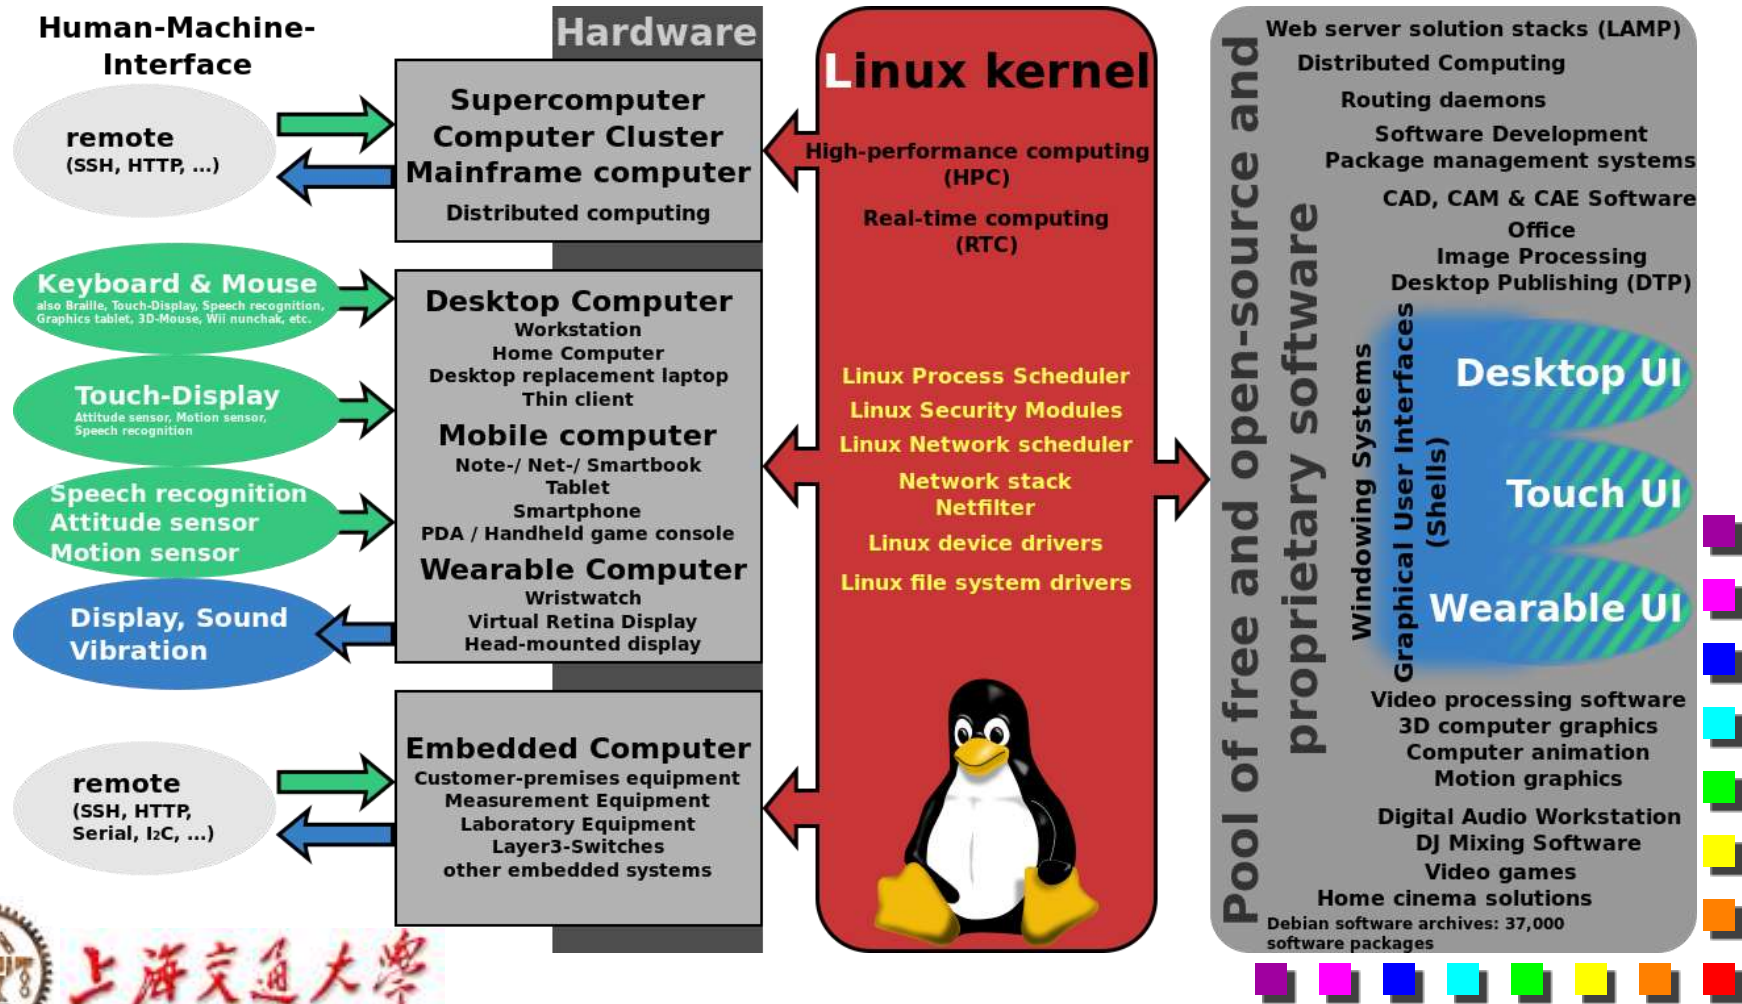


Board View of Linux Kernel


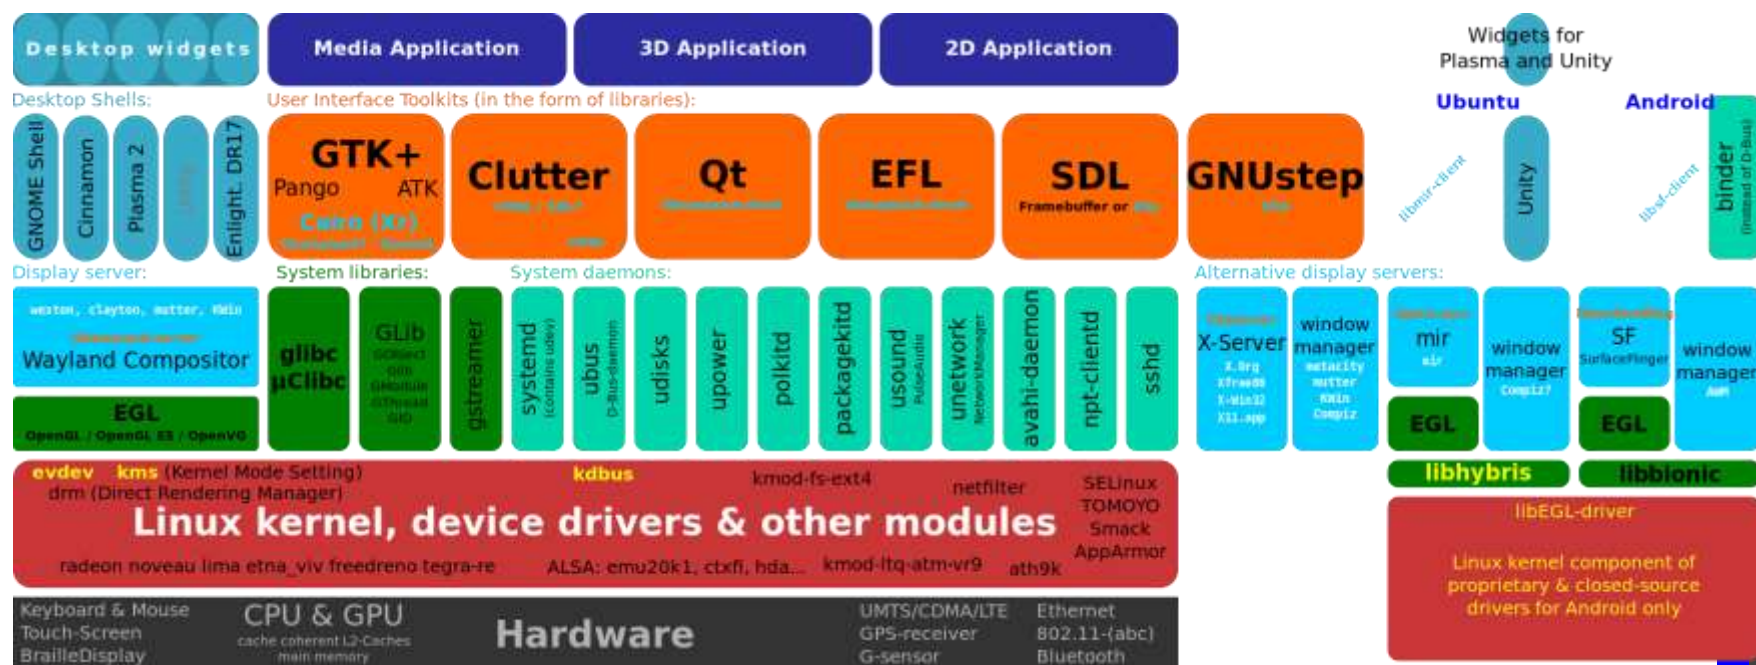


上海交通大学

Various Types of Hardware Are Running on Linux



Software Components of Linux Desktop




上海交通大学

- 0.01 -- 17 Sep. 1991
- 1.0 -- 14 Mar. 1994
- 2.0 -- 9 June 1996
- 2.4 -- 4 Jan. 2001
- 2.6 -- 18 Dec. 2003
- 3.0 -- 21 July 2011
- 4.0 -- 12 Apr. 2015
- 4.5 -- 21 Feb. 2016



上海交通大学

Lines of Source Code in Linux



kernel.org

The Linux Kernel Archives



[About](#) [Contact us](#) [FAQ](#) [Releases](#) [Signatures](#) [Site news](#)

Protocol	Location
HTTP	https://www.kernel.org/pub/
GIT	https://git.kernel.org/
RSYNC	rsync://rsync.kernel.org/pub/

Latest Stable Kernel:



4.4.2


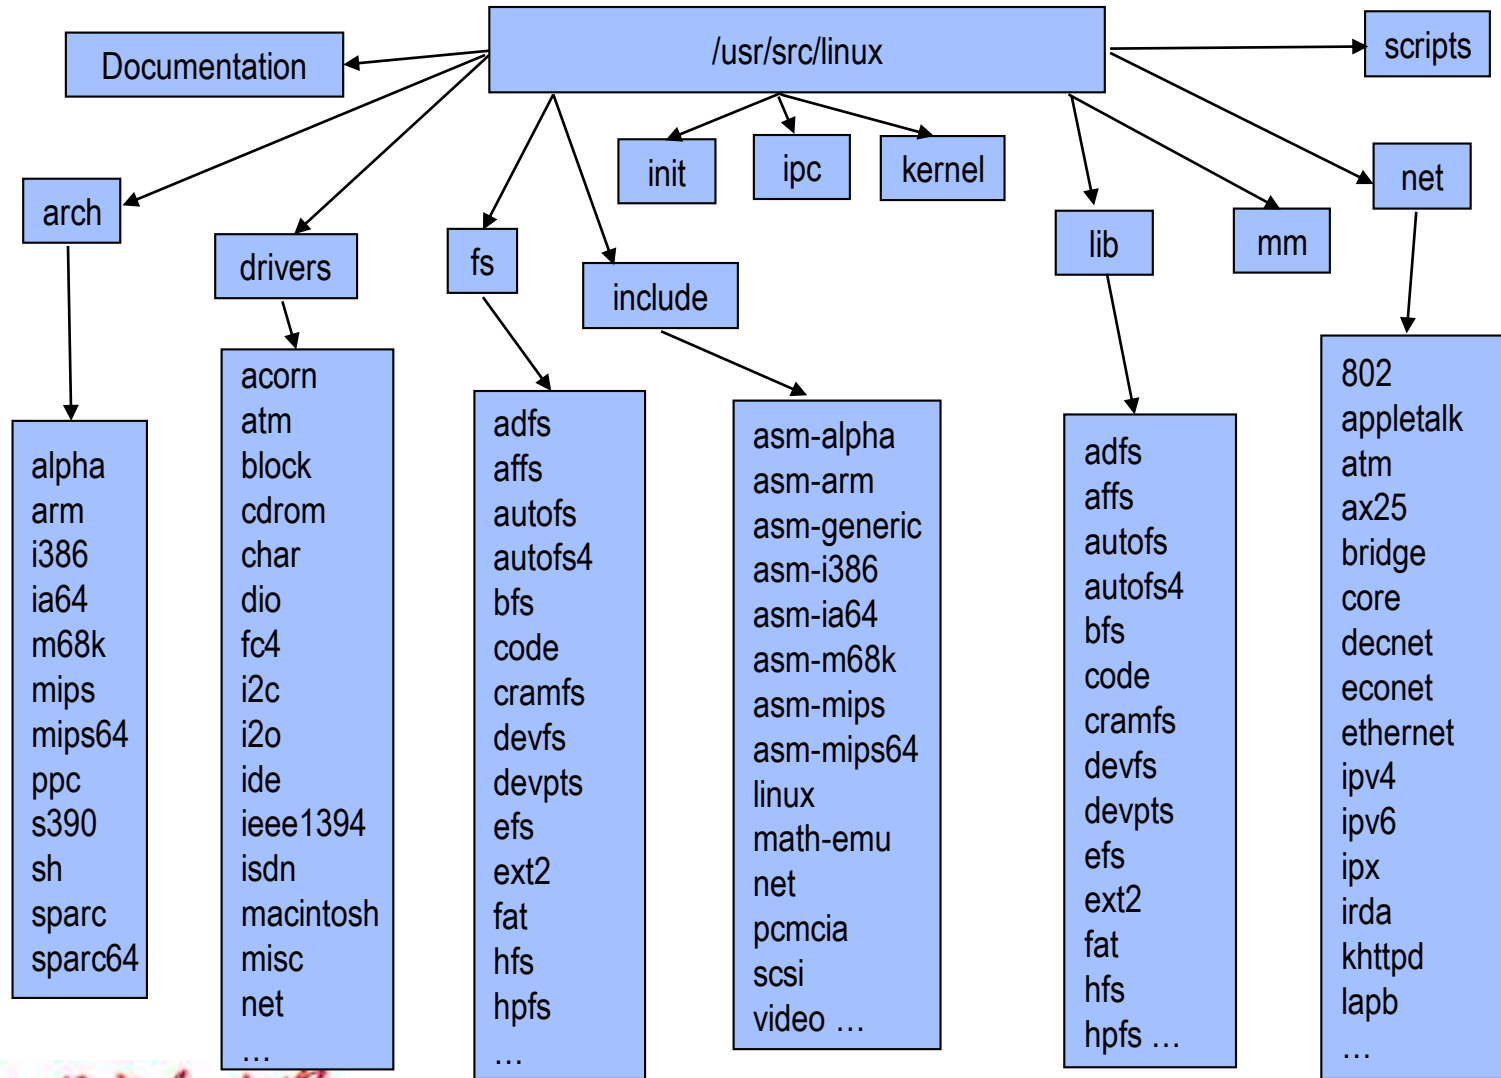
mainline:	4.5-rc5	2016-02-20	[tar.xz]	[pgp]	[patch]	[view diff]	[browse]	
stable:	4.4.2	2016-02-17	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
stable:	4.3.6 [EOL]	2016-02-19	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	4.1.18	2016-02-15	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	3.18.27	2016-02-15	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	3.14.61	2016-02-17	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	3.12.54	2016-02-15	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	3.10.97	2016-02-19	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	3.4.110	2015-10-22	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	3.2.77	2016-02-13	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	2.6.32.70	2016-01-29	[tar.xz]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
linux-next:	next-20160225	2016-02-25						[browse]



上海交通大学

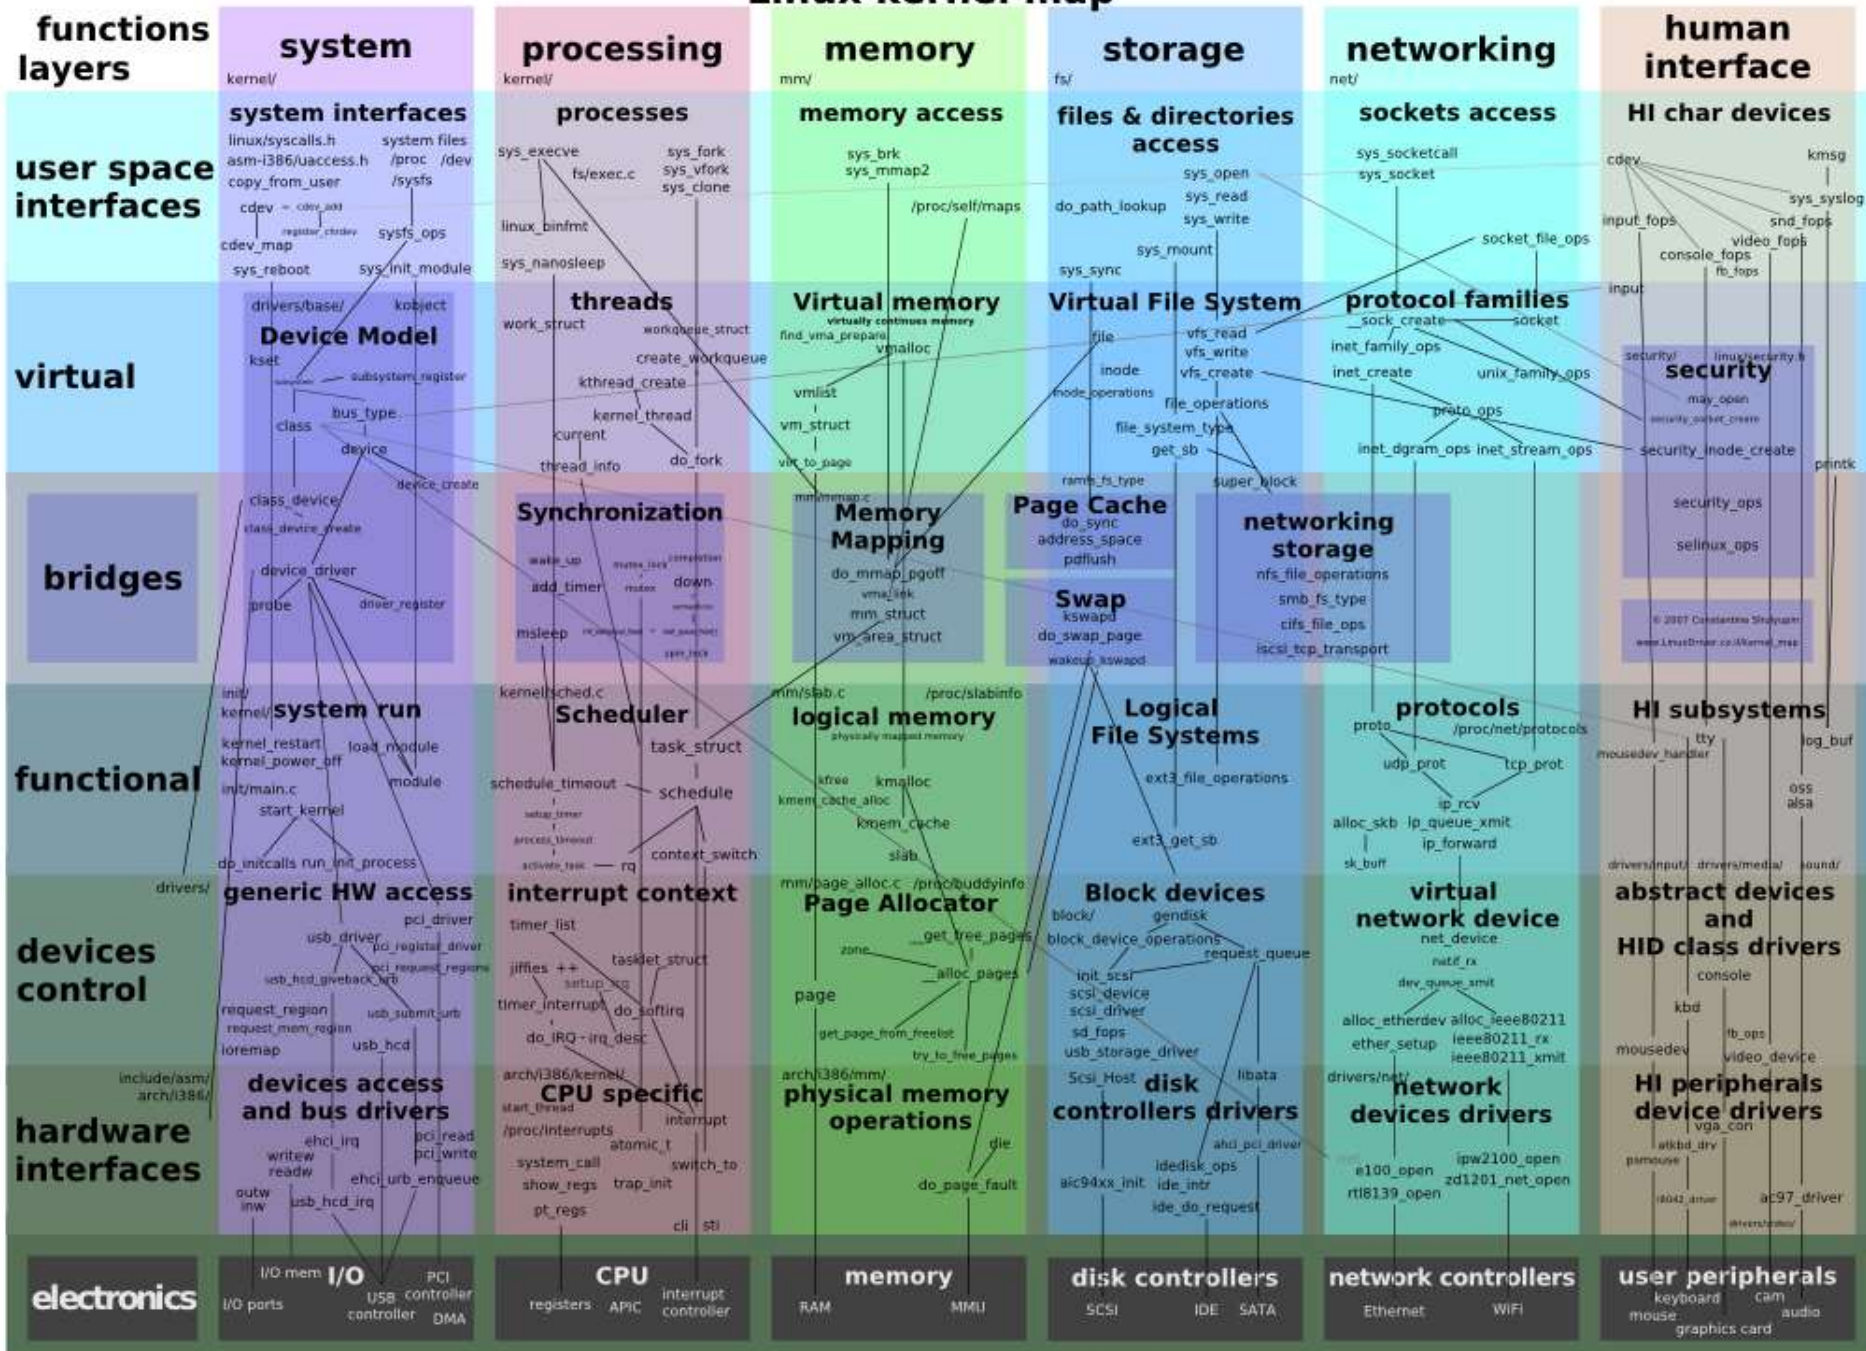


Linux Source Tree Layout (2.6 or earlier)



上海交通大学

Linux kernel map



Linux Source Tree (1)

- All directories are grouped under the root entry "/"
- **root** - The home directory for the root user
- **home** - Contains the user's home directories along with directories for services
 - ftp
 - HTTP
 - samba



Linux Source Tree (2)

- **bin** - Commands needed during booting up that might be needed by normal users
- **sbin** - Like bin but commands are not intended for normal users. Commands run by LINUX.
- **proc** - This filesystem is not on a disk. It is a virtual filesystem that exists in the kernels imagination which is memory
 - 1 - A directory with info about process number 1. Each process has a directory below proc.



Linux Source Tree (3)

- **usr** - Contains all commands, libraries, man pages, games and static files for normal operation.
 - **bin** - Almost all user commands. some commands are in /bin or /usr/local/bin.
 - **sbin** - System admin commands not needed on the root filesystem. e.g., most server programs.
 - **include** - Header files for the C programming language. Should be below /usr/lib for consistency.
 - **lib** - Unchanging data files for programs and subsystems
 - **local** - The place for locally installed software and other files.
 - **man** - Manual pages
 - **info** - Info documents
 - **doc** - Documentation
 - **tmp**
 - **X11R6** - The X windows system files. There is a directory similar to usr below this directory.
 - **X386** - Like X11R6 but for X11 release 5



Linux Source Tree (4)

- **boot** - Files used by the bootstrap loader. Kernel images are often kept here.
- **lib** - Shared libraries needed by the programs on the root filesystem
- **modules** - Loadable kernel modules, especially those needed to boot the system after disasters.
- **dev** - Device files
- **etc** - Configuration files specific to the machine.
- **skel** - When a home directory is created it is initialized with files from this directory
- **sysconfig** - Files that configure the linux system for devices.



Linux Source Tree (5)

- **var** - Contains files that change for mail, news, printers log files, man pages, temp files
 - **file**
 - **lib** - Files that change while the system is running normally
 - **local** - Variable data for programs installed in /usr/local.
 - **lock** - Lock files. Used by a program to indicate it is using a particular device or file
 - **log** - Log files from programs such as login and syslog which logs all logins and logouts.
 - **run** - Files that contain information about the system that is valid until the system is next booted
 - **spool** - Directories for mail, printer spools, news and other spooled work.
 - **tmp** - Temporary files that are large or need to exist for longer than they should in /tmp.
 - **catman** - A cache for man pages that are formatted on demand




Linux Source Tree (6)

- **mnt** - Mount points for temporary mounts by the system administrator.
- **tmp** - Temporary files. Programs running after bootup should use /var/tmp



linux/arch

- Subdirectories for each current port.
- Each contains **kernel**, **lib**, **mm**, **boot** and other directories whose contents override code stubs in architecture independent code.
- **lib** contains highly-optimized common utility routines such as memcpy, checksums, etc.
- **arch** as of 2.4:
 - alpha, arm, i386, ia64, m68k, mips, mips64.
 - ppc, s390, sh, sparc, sparc64.



linux/drivers

- Largest amount of code in the kernel tree (~1.5M).
- device, bus, platform and general directories.
- drivers/char – n_tty.c is the default line discipline.
- drivers/block – elevator.c, genhd.c, linear.c, ll_rw_blk.c, raidN.c.
- drivers/net – specific drivers and general routines Space.c and net_init.c.
- drivers/scsi – scsi_*.c files are generic; sd.c (disk), sr.c (CD-ROM), st.c (tape), sg.c (generic).
- General:
 - cdrom, ide, isdn, parport, pcmcia, pnp, sound, telephony, video.
- Buses – fc4, i2c, nubus, pci, sbus, tc, usb.
- Platforms – acorn, macintosh, s390, sgi.



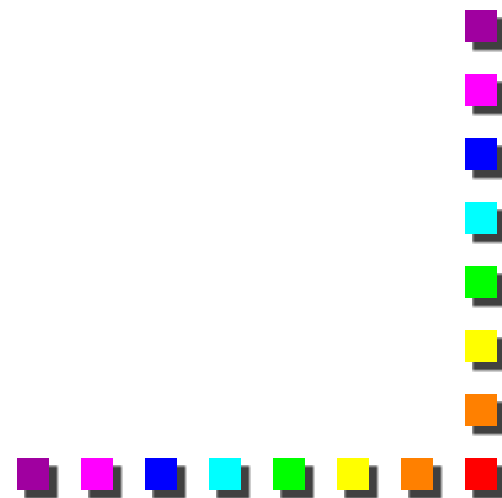
linux/fs

- Contains:
 - virtual filesystem (VFS) framework.
 - subdirectories for actual filesystems.
- vfs-related files:
 - exec.c, binfmt_*.c - files for mapping new process images.
 - devices.c, blk_dev.c – device registration, block device support.
 - super.c, filesystems.c.
 - inode.c, dcache.c, namei.c, buffer.c, file_table.c.
 - open.c, read_write.c, select.c, pipe.c, fifo.c.
 - fcntl.c, ioctl.c, locks.c, dquot.c, stat.c.



linux/include

- include/asm-*:
 - Architecture-dependent include subdirectories.
- include/linux:
 - Header info needed both by the kernel and user apps.
 - Usually linked to /usr/include/linux.
 - Kernel-only portions guarded by #ifdefs
 - #ifdef __KERNEL__
 - /* kernel stuff */
 - #endif
- Other directories:
 - math-emu, net, pcmcia, scsi, video.




linux/init

- Just two files: version.c, main.c.
- version.c – contains the version banner that prints at boot.
- main.c – architecture-independent boot code.
- start_kernel is the primary entry point.



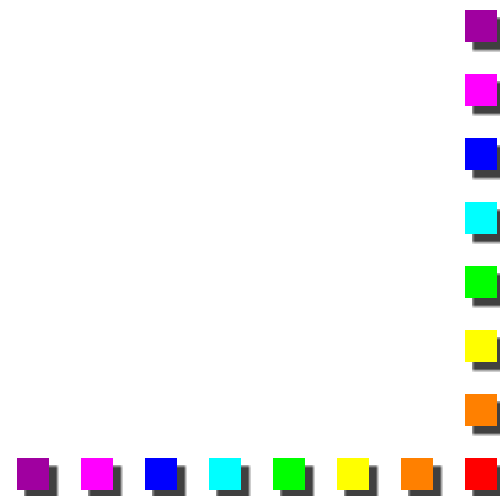
linux/ipc

- System V IPC facilities.
- If disabled at compile-time, util.c exports stubs that simply return `–ENOSYS`.
- One file for each facility:
 - sem.c – semaphores.
 - shm.c – shared memory.
 - msg.c – message queues.



linux/kernel

- The core kernel code.
- sched.c – “the main kernel file”:
 - scheduler, wait queues, timers, alarms, task queues.
- Process control:
 - fork.c, exec.c, signal.c, exit.c etc...
- Kernel module support:
 - kmod.c, ksyms.c, module.c.
- Other operations:
 - time.c, resource.c, dma.c, softirq.c, itimer.c.
 - printk.c, info.c, panic.c, sysctl.c, sys.c.




linux/lib

- kernel code cannot call standard C library routines.
- Files:
 - brlock.c – “Big Reader” spinlocks.
 - cmdline.c – kernel command line parsing routines.
 - errno.c – global definition of errno.
 - inflate.c – “gunzip” part of gzip.c used during boot.
 - string.c – portable string code.
 - Usually replaced by optimized, architecture-dependent routines.
 - vsprintf.c – libc replacement.




linux/mm

- Paging and swapping:
 - swap.c, swapfile.c (paging devices), swap_state.c (cache).
 - vmscan.c – paging policies, kswapd.
 - page_io.c – low-level page transfer.
- Allocation and deallocation:
 - slab.c – slab allocator.
 - page_alloc.c – page-based allocator.
 - vmalloc.c – kernel virtual-memory allocator.
- Memory mapping:
 - memory.c – paging, fault-handling, page table code.
 - filemap.c – file mapping.
 - mmap.c, mremap.c, mlock.c, mprotect.c.



linux/scripts

- Scripts for:
 - Menu-based kernel configuration.
 - Kernel patching.
 - Generating kernel documentation.



Summary

- Linux is a modular, UNIX-like monolithic kernel.
- Kernel is the heart of the OS that executes with special hardware permission (kernel mode).
- “Core kernel” provides framework, data structures, support for drivers, modules, subsystems.
- Architecture dependent source sub-trees live in /arch.



Project 1: Compile the Linux Kernel



上海交通大學

Preparation



- For each student
 - Prepare one laptop for each group
 - Install Virtual Machine Software: **VMware Workstation** or **VirtualBox**
 - Install the Linux Operating System in a virtual machine: **RedHat Enterprise Linux (RHEL)**, **Fedora**, Gentoo, **Ubuntu**, etc.
 - Linux kernel version: 4.3 or earlier



上海交通大学

Compile The Linux Kernel (1)

- Download the latest kernel version (4.4.2, updated on 2016-02-17) from <http://www.kernel.org>
 - Download the complete source code (not a patch or change log)
- Copy to a directory
 - **4.X.X is the kernel version number**

```
# cp linux-4.X.X.tar.gz /usr/src
```

```
# cd /usr/src
```

- Extract the kernel

```
# tar -xzvf linux-4.X.X.tar.gz or
```

```
# tar -xjvf linux-4.X.X.tar.bz2
```

上海交通大学



Compile The Linux Kernel (2)

- Preparation

cd /usr/src/linux

make clean

make mrproper

Note: make clean and make mrproper can be passed with the first compilation.



上海交通大學



Compile The Linux Kernel (3)

- Configure the kernel

make menuconfig

/*creates a menu where you can browse options on what the kernel supports*/

make xconfig

/*same as menuconfig, except that now the configuration menu is graphics based*/


make oldconfig

/*minor revision on previous kernel*/

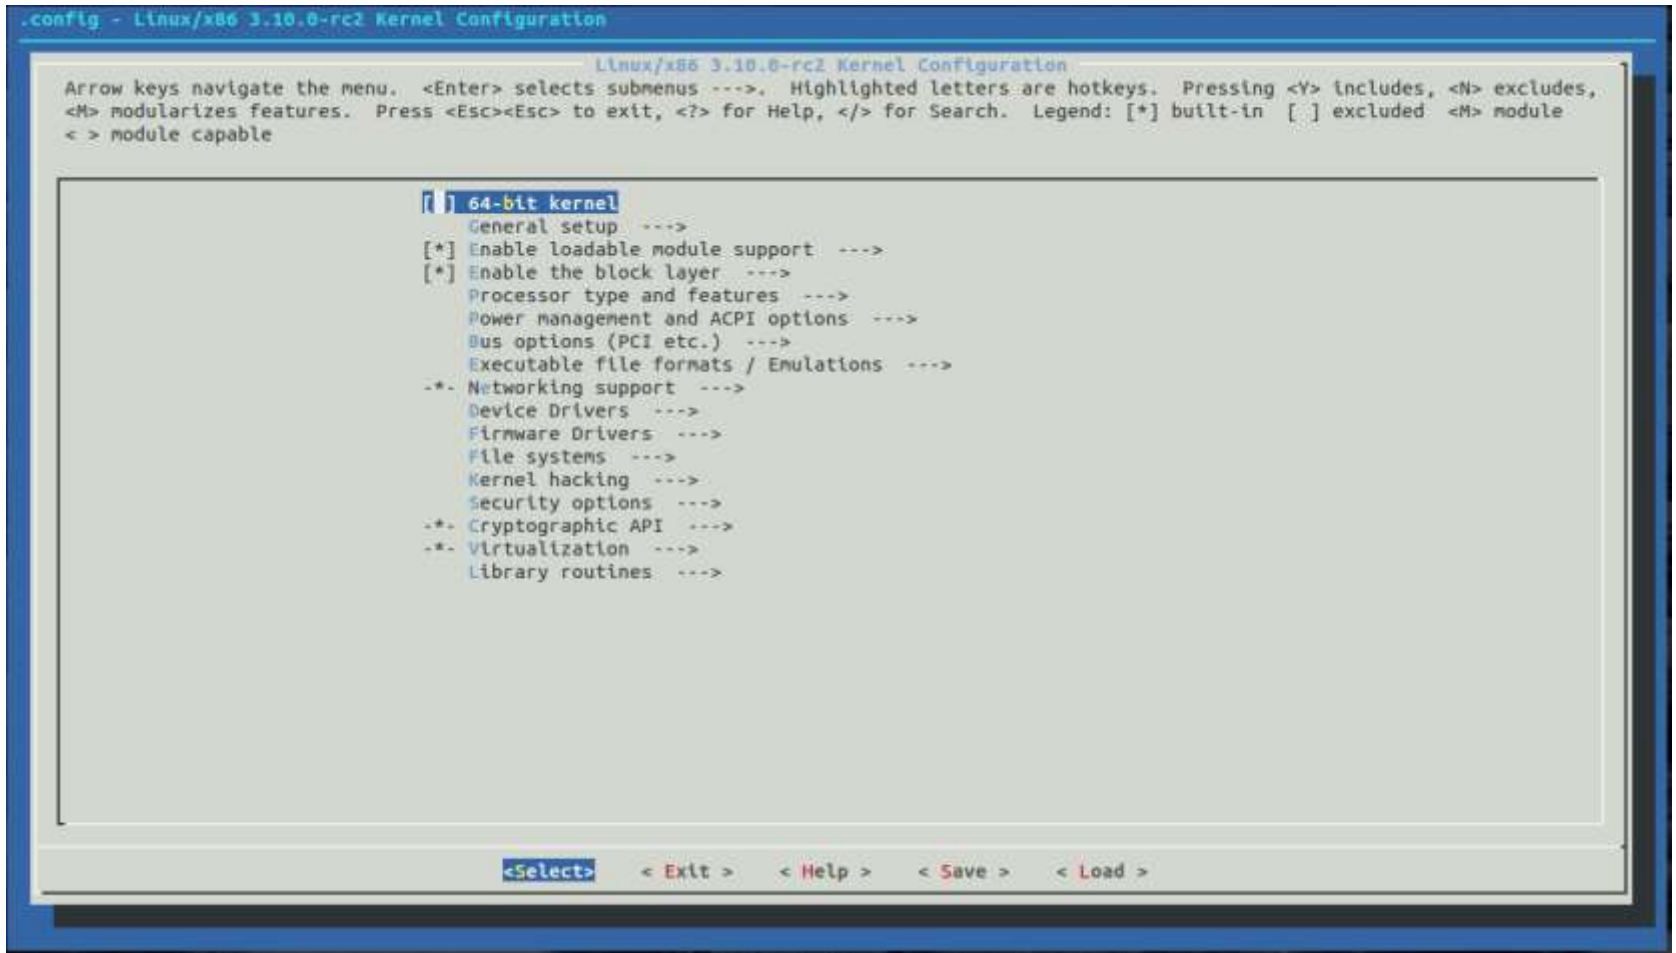
make config **/*not be recommended*/**




上海交通大學



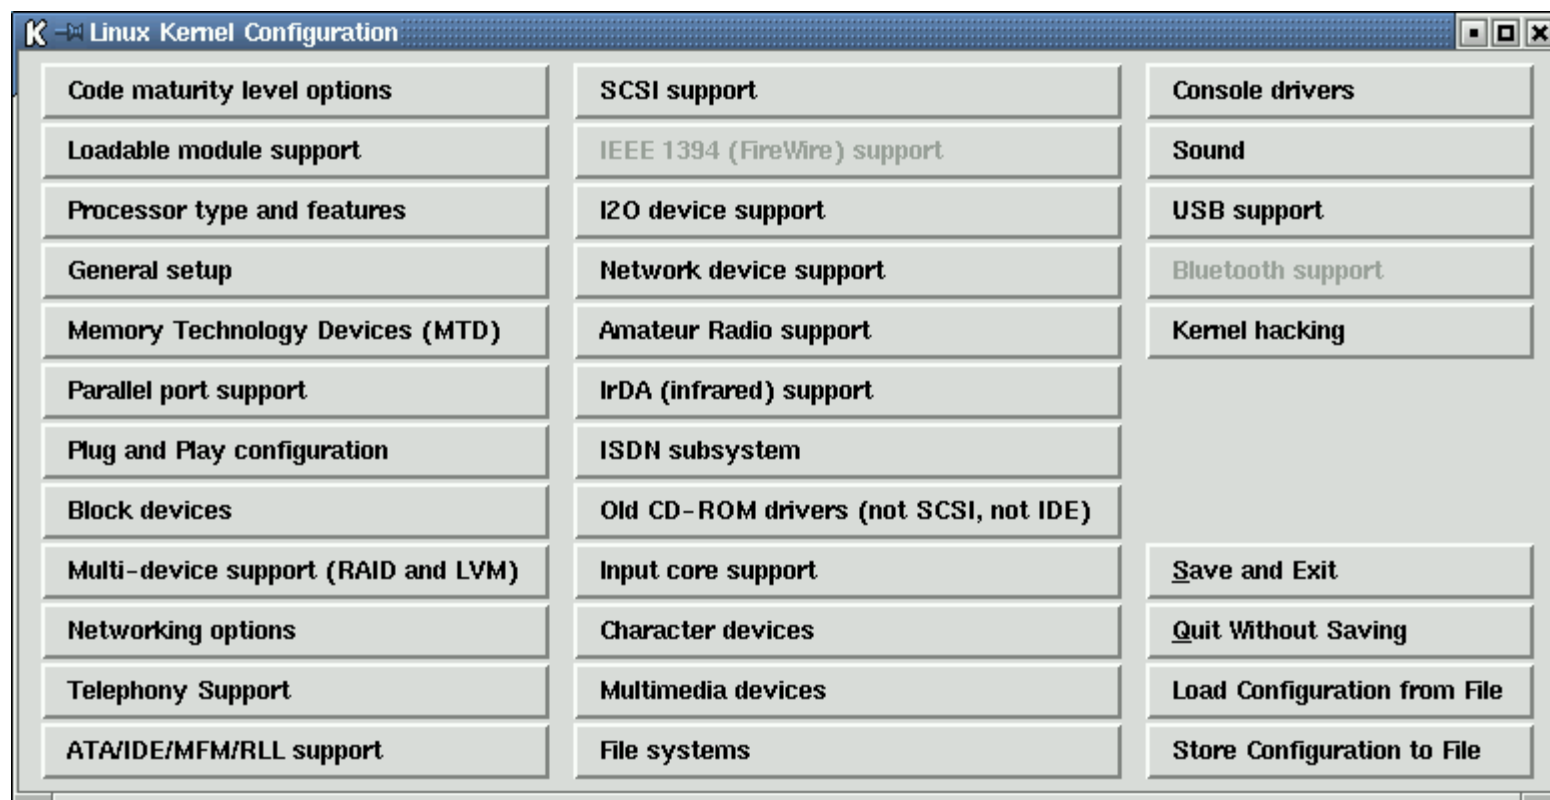
Make menuconfig




上海交通大学



Make Xconfig



上海交通大学



Compile The Linux Kernel (4)

- Customize the Kernel
 - e.g., you may add support for NTFS file system from “File System >> DOS/FAT/NT/ >> select NTFS file system support”



Compile The Linux Kernel (5)

- Compile and Install the Kernel (~20 min)

make

make modules_install

make install

- Update GRUB (or LILO)

cd /boot

mkinitrd -o initrd.img-<kernel version>

- Reboot the system



上海交通大學

