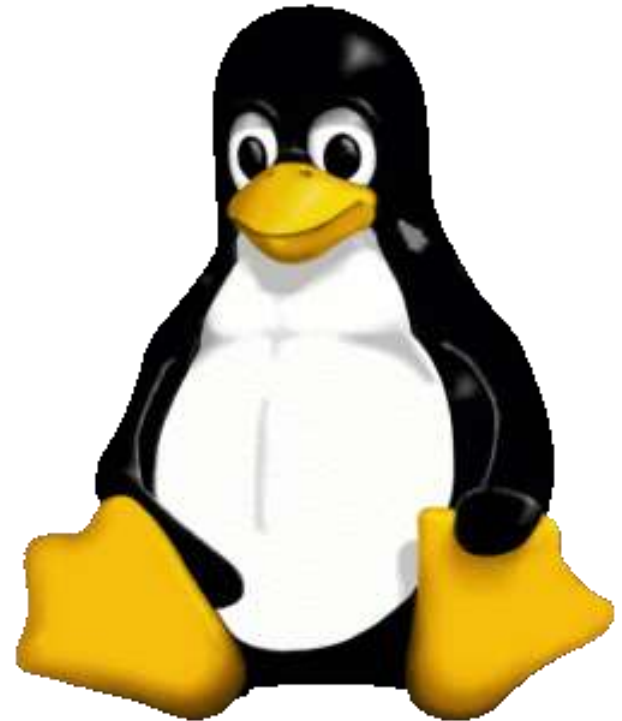


# CS353 Linux Kernel

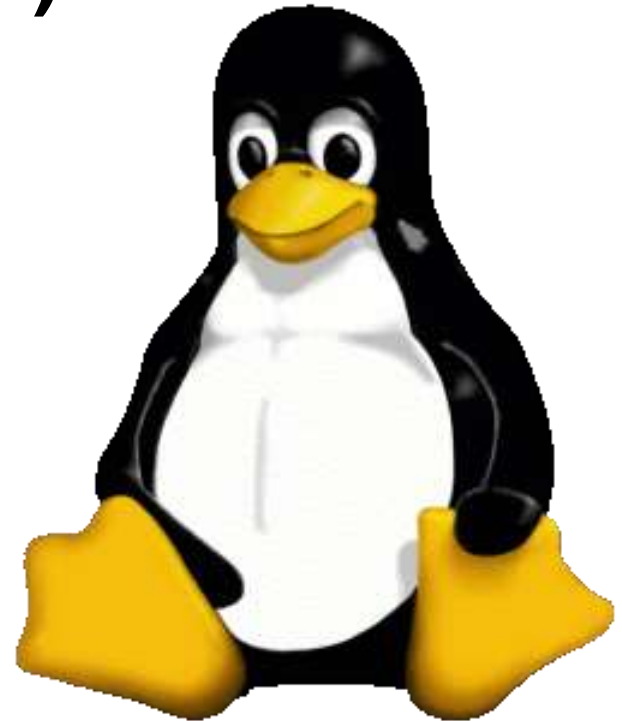
**Chentao Wu 吴晨涛**  
**Associate Professor**  
**Dept. of CSE, SJTU**  
**wuct@cs.sjtu.edu.cn**



上海交通大学

# 7C. Advanced File System Technology (Object-based Storage)

Chentao Wu  
Associate Professor  
Dept. of CSE, SJTU  
[wuct@cs.sjtu.edu.cn](mailto:wuct@cs.sjtu.edu.cn)



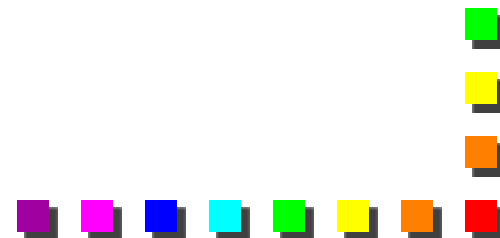
上海交通大學



# Object Storage Technology

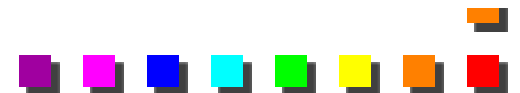


上海交通大學



# Agenda

- Object storage background and history
  - ◆ Abstract data containers
  - ◆ NASD, OSD, HTTP
- Objects and File Systems
  - ◆ Lustre, PanFS, Ceph, many others
- Objects and Web Storage
  - ◆ S3, Azure, Swift, many others



# Two Paths to Objects

## ➤ Storage Devices

- ◆ Move smarts into the device (NASD)
  - › Network Attached Secure Disk
- ◆ Raise the level of abstraction
  - › Containers
  - › Attributes
  - › Security
- ◆ SCSI Model
  - › OSD command set

## ➤ Web Services

- ◆ Add storage abstraction to a web-based system
  - › Containers
  - › Metadata
  - › Security
- ◆ REST Model
  - › HTTP protocol

# Some History



## ❖ NASD (Network Attached Secure Disk)

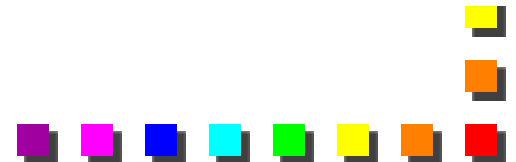
- ◆ 1990's research by Dr. Garth Gibson about moving intelligence into the storage device
- ◆ Google cites NASD as inspiration for data node in its file system

## ❖ OSD (Object-based Storage Device)

- ◆ Standards effort in the early 2000's created a SCSI command set for objects
- ◆ There is a storage device behind this interface

## ❖ HTTP (HyperText Transfer Protocol)

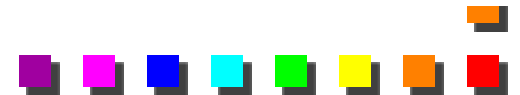
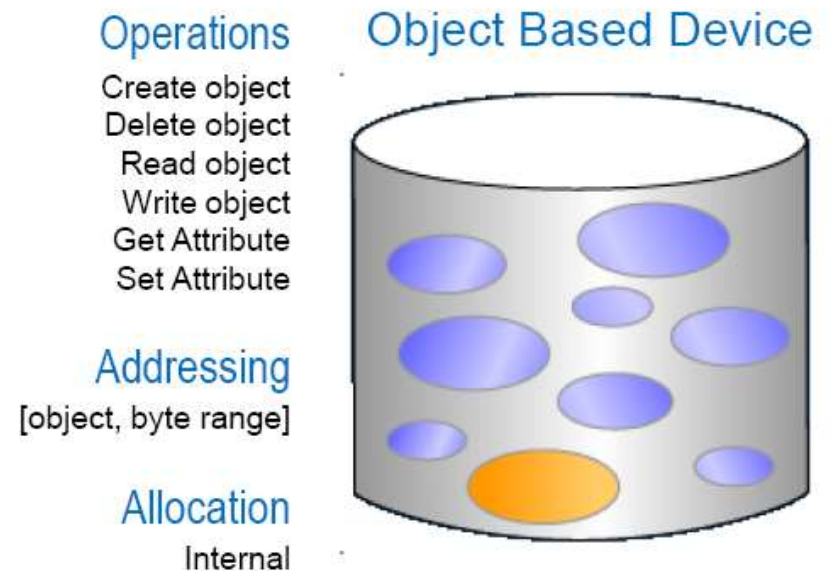
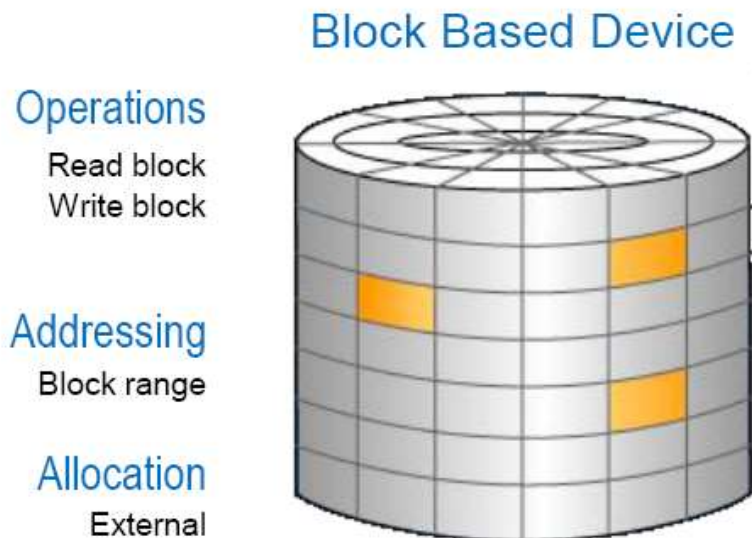
- ◆ A simple put/get protocol for the world-wide web
- ◆ There is an arbitrary service behind this interface



# Object Storage

## ► Objects are containers for data and attributes

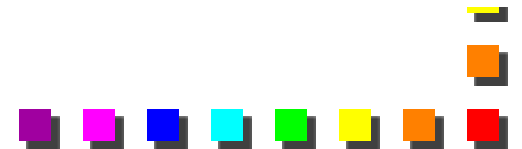
- ◆ Every file system has an *inode* that is data blocks plus attributes
- ◆ They are created, deleted, read, written, and have attributes





# Object Storage Device (OSD)

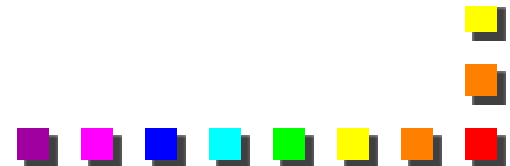
- OSDs hold objects, which are like files in a simple file system
  - ◆ Objects are dynamically created and freed
  - ◆ Object are variable length
  - ◆ Objects have extensible attributes
  - ◆ Objects are identified by a 64 bit Object ID (OID)
  - ◆ Objects in an OSD are grouped within partitions, which are identified by a 64 bit Partition ID
    - 64 bit OID plus 64 bit PID gives a 128 bit namespace
- OSDs manage space allocation of Objects
  - ◆ A 4TB disk has 1 billion 4KB blocks – OSD hides this





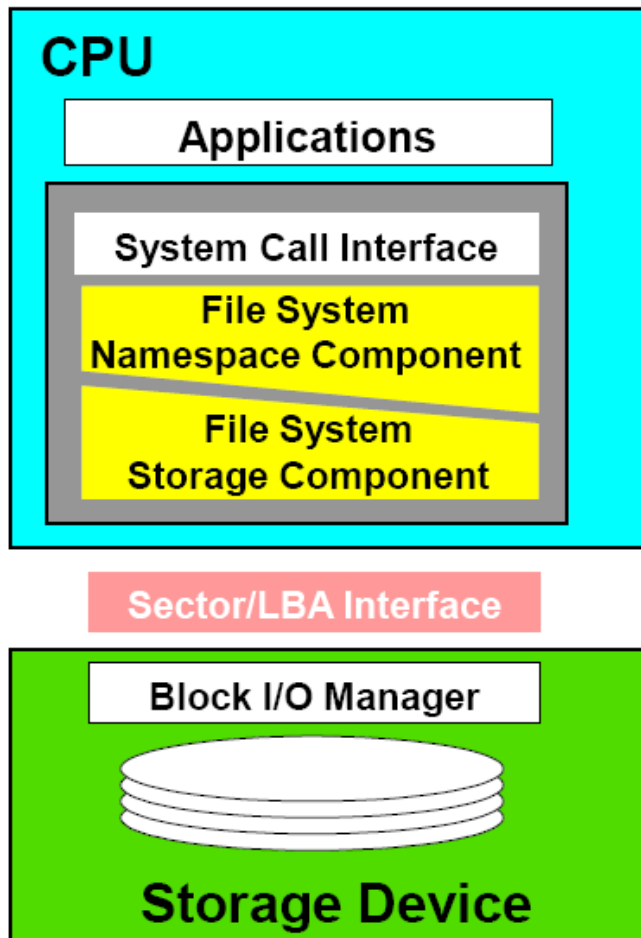
# OSD Standards

- There is a standard for OSDs under ANSI INCITS T10 (the SCSI specification)
  - ◆ *ANSI INCITS 458-2011*
  - ◆ OSD-1 is basic functionality
    - › Read, write, create, delete objects and partitions
    - › Security model, Capabilities, manage shared secrets and working keys
  - ◆ OSD-2 adds:
    - › Snapshots
    - › Collections of Objects
    - › Extended exception handling and recovery

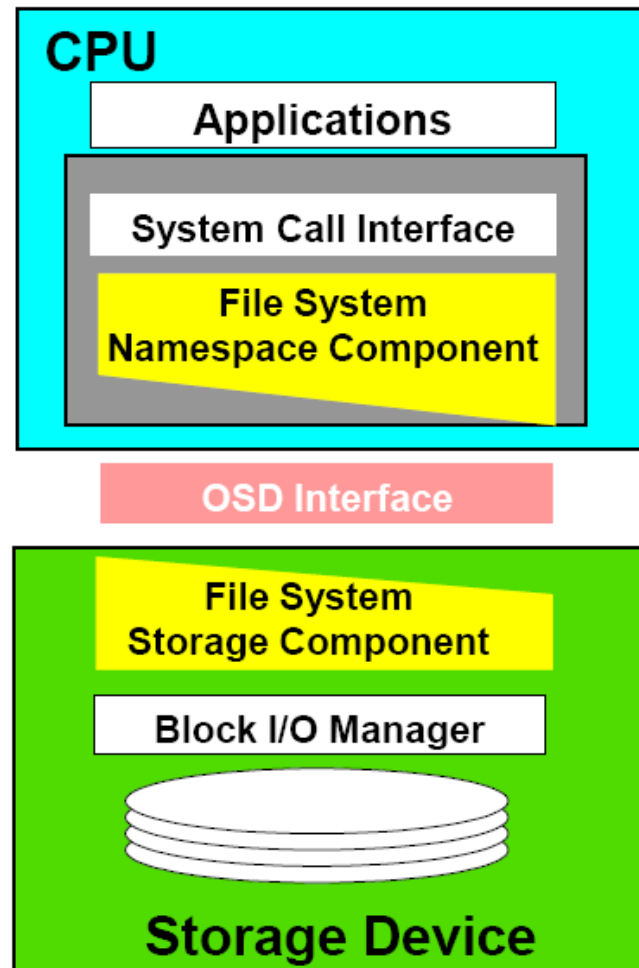


# Objects and the File System Stack

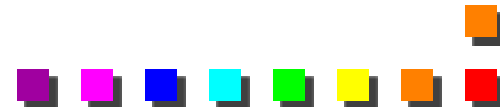
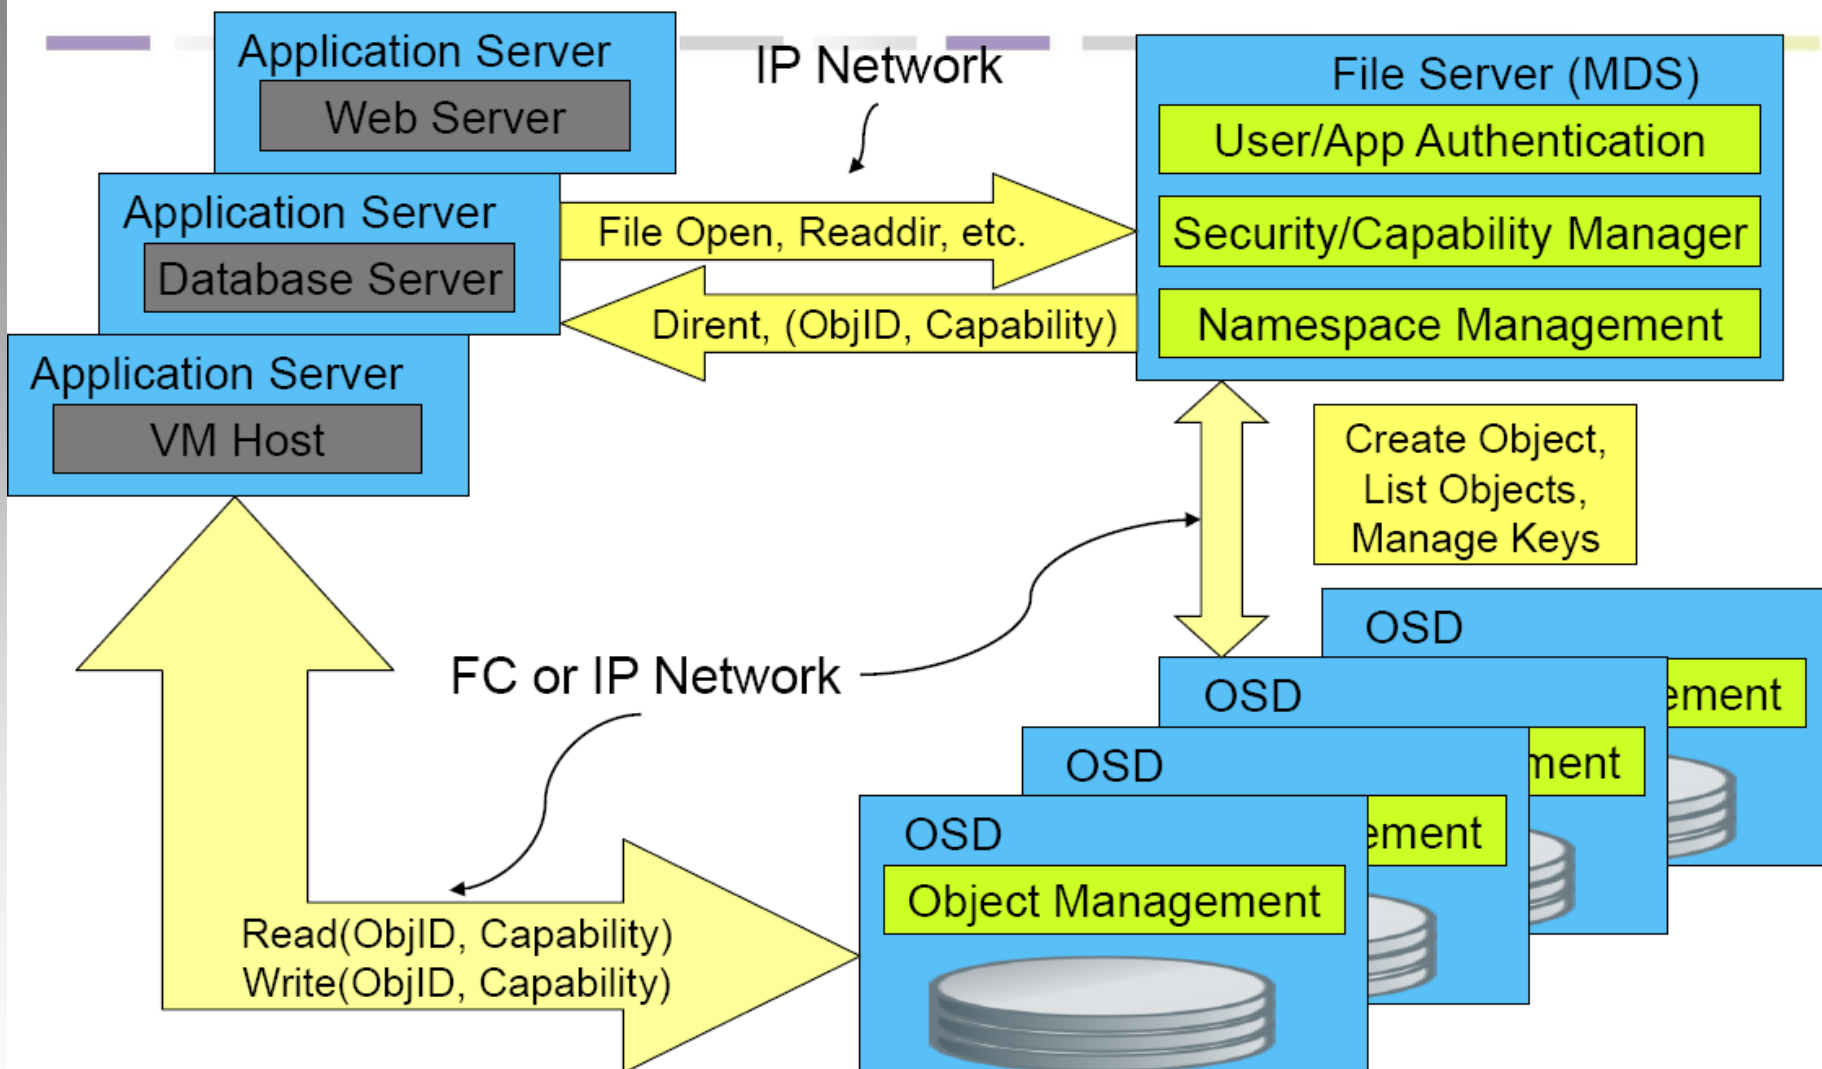
Traditional File System



Object-based File System

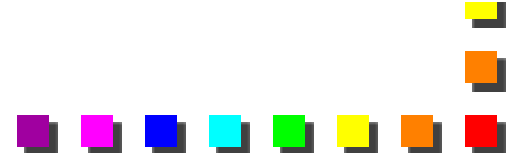


# Scaling OSD File Systems



# Why Objects Help Scaling

- 90% of File System cycles are in the read/write path
  - ◆ Block allocation is expensive
  - ◆ Data transfer is expensive
  - ◆ OSD offloads both of these from the file server
  - ◆ Security model allows direct access from clients
- Higher level interfaces allow optimization
  - ◆ The more function behind an API, the less often you have to use the API to get your work done
- Higher level interfaces provide more semantics
  - ◆ User authentication and access control
  - ◆ Namespaces and indexing

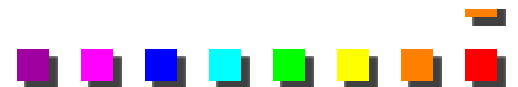


# OSD Capabilities

- Unlike disks, where access is granted on an all or nothing basis, OSDs grant or deny access to individual objects based on *Capabilities*
- A Capability must accompany each request to read or write an object
  - ◆ Capabilities are cryptographically signed by the Security Manager and verified (and enforced) by the OSD
  - ◆ A Capability to access an object is created by the Security Manager, and given to the client (application server) accessing the object
  - ◆ Capabilities can be revoked by changing an attribute on the object

# OSD Security Model

- OSD and File Server know a secret key
  - ◆ Working keys are periodically generated from a master key
- File server authenticates clients and makes access control policy decisions
  - ◆ Access decision is captured in a capability that is signed with the secret key
  - ◆ Capability identifies object, expire time, allowed operations, etc.
- Client signs requests using the capability signature as a signing key
  - ◆ OSD verifies the signature before allowing access
  - ◆ OSD doesn't know about users, ACLs, or whatever policy mechanism the File Server is using



# Object Storage File Systems

## ➤ Lustre

- ◆ Custom OSS/OST model
- ◆ Single metadata server

## ➤ PanFS

- ◆ ANSI T10 OSD model
- ◆ Multiple metadata servers

## ➤ Ceph

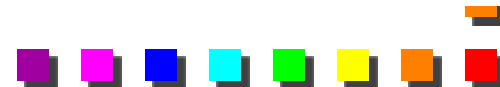
- ◆ Custom OSD model
- ◆ Crush metadata distribution

## ➤ pNFS

- ◆ Out-of-band metadata service for NFSv4.1
- ◆ T10 Objects, Files, Blocks as data services

## ➤ These systems scale

- ◆ 1000's of disks (i.e., PB's)
- ◆ 1000's of clients
- ◆ 100's GB/sec
- ◆ all in one file system

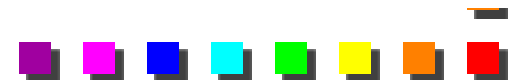
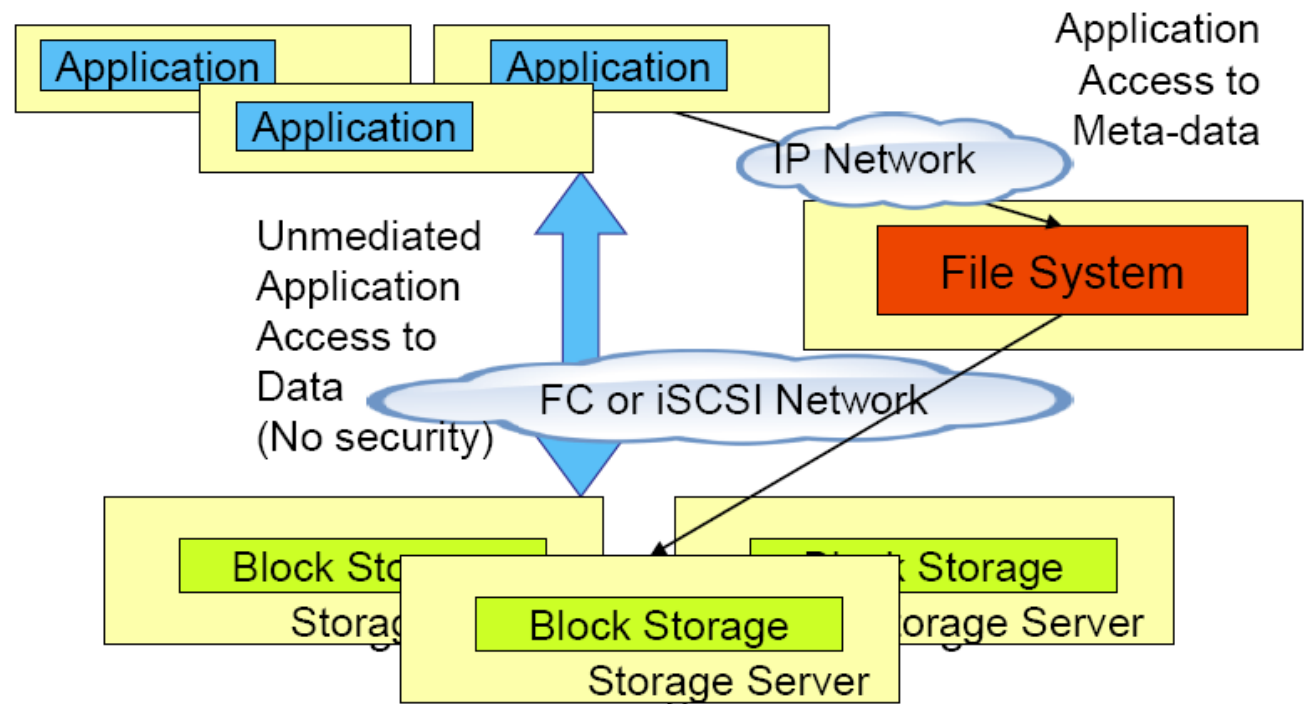




# What about SAN File Systems?

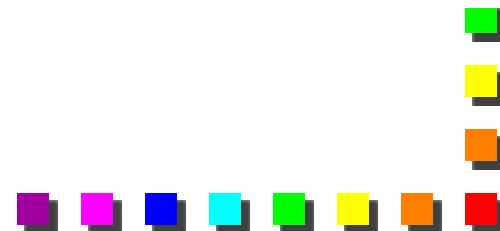
## ◆ SAN file systems out-of-band metadata service

- ◆ Security model does not support fine grain sharing
- ◆ Block allocation must be managed by the file system
- ◆ This can work well at small scale (many products)



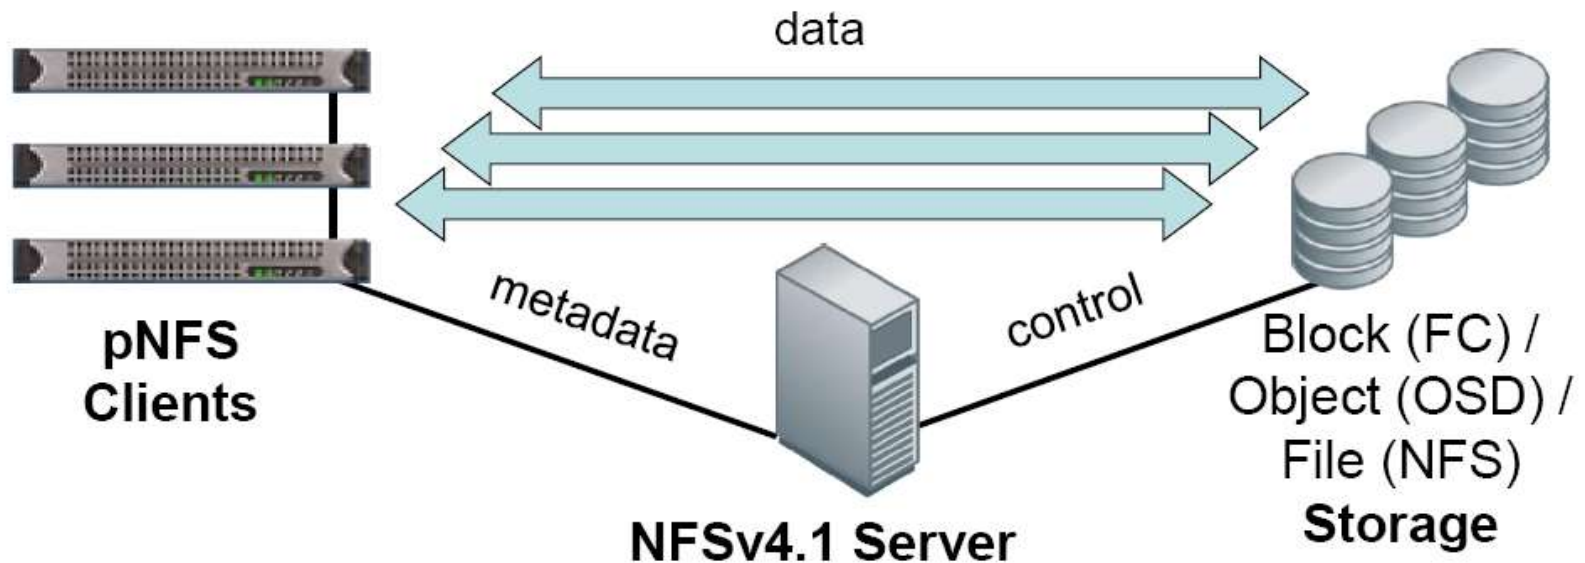
# Why a Standard for Parallel I/O?

- NFS is the only network file system standard
  - ◆ IETF Standard
  - ◆ Proprietary file systems have unique advantages, but aren't right for everyone
    - PanFS, Lustre, GPFS, IBRIX, CXFS, HDFS, etc.
- pNFS widens the playing field
  - ◆ Most major NFS vendors have announced pNFS support
  - ◆ Broader market benefits vendors
  - ◆ More competition benefits customers
- NFS standard effort very active with 4.2 and beyond



# Parallel IO in the NFS Standard

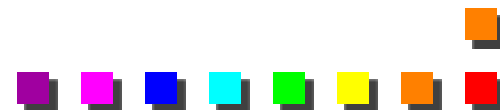
- ◆ How do you get the NFS file server out of the data path?
  - › Add *Layouts*, and introduce a level of indirection
  - › Conceptually small addition to the standard ☺
- ◆ What is the benefit?
  - › Better overall performance from a large storage system
    - Load balancing among storage devices
  - › Standard client for high performance, large scale storage systems



# pNFS Status

---

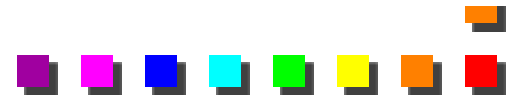
- ◆ RFCs for NFSv4.1, pNFS-objects, and pNFS-blocks published January 2010
  - ◆ RFC 5661 - Network File System (NFS) Version 4 Minor Version 1 Protocol
  - ◆ RFC 5662 - Network File System (NFS) Version 4 Minor Version 1 External Data Representation Standard (XDR) Description
  - ◆ RFC 5663 - Parallel NFS (pNFS) Block/Volume Layout
  - ◆ RFC 5664 - Object-Based Parallel NFS (pNFS) Operations
- ◆ Linux 3.4 kernel (mid 2012) has all of the patches necessary for the pNFS client
  - ◆ RHEL 6.4 and SLES offer support via back-ports
  - ◆ Enterprise distros based on more modern kernels by end of 2013



# Web Object Features

---

- REST-ful API (i.e., web-based)
- Security/Authentication tied to Billing
- Metadata capabilities
- Highly available
- Loosely consistent
- Data Storage
  - ◆ Blobs
  - ◆ Tables
  - ◆ Queues
- Other related APIs (compute, search, etc.)
  - ◆ Storage API is relatively simple in comparison

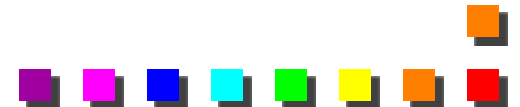


# RESTful API

## ◆ Representational State Transfer (i.e., it's the web)

- ◆ Client-server
- ◆ Stateless
  - › Every request transmits session state (i.e., Cookies)
- ◆ Cacheable (or not)
  - › Explicit in the protocol
- ◆ Layered
  - › transparent proxies
- ◆ Code on demand (optional)
  - › javascript
- ◆ Uniform
  - › URI, Content-Type, MIME headers, ...

Simple, flexible,  
text-based protocol  
that is easy to  
implement and  
extend.  
It is just HTTP



# Simple HTTP example

% telnet www.google.com 80

GET /index.html HTTP/1.0

(blank line)

HTTP/1.0 200 OK

Date: Wed, 13 Feb 2013 07:24:07 GMT

Content-Type: text/html; charset=ISO-8859-1

<html>

<head><title>Google</title></head>

<body><img src= /images/srpr/logo3w.png>

<form><input type=text name=q>

<input type=submit value="Google Search" name="search">

<input type=submit value="I'm Feeling Lucky" name="lucky">

</form></body></html>

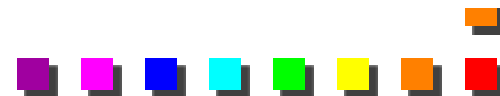
GET  
parameters  
metadata



REPLY  
metadata  
data



上海交通大學





# HTTP and Objects

## ➤ Request specifies method and object:

- ◆ Operation: GET, POST, PUT, HEAD, COPY
- ◆ Object ID (/index.html)

This is a method  
call on an object

## ➤ Parameters use MIME format borrowed from email

- ◆ Content-type: utf8;
- ◆ Set-Cookie: tracking=1234567;

These are  
parameters

## ➤ And a data payload

- ◆ Optional
- ◆ Separated from parameters with a blank line (like email)

This is data

## ➤ Response has identical structure

- ◆ Status line, key-value parameters, optional data payload

# Security

## ➤ Shared Keys and Signatures

- ◆ Shared secret is typical
- ◆ Obtained from service provider
- ◆ Two factor, or public key/private key also possible

## ➤ Signature computed over the request

- ◆ GET/PUT line and some/all the request metadata
- ◆ Signed with the secret using SHA256
- ◆ Signature appears as metadata on the request

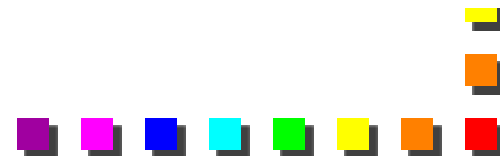
**Authorization: SharedKey**

**myaccount:ctzMq410TV3wS7upTBcunJTDLEJwMAZuFPfr0mrrA08=**

## ➤ Details handled by an SDK

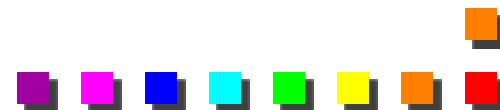


上海交通大學



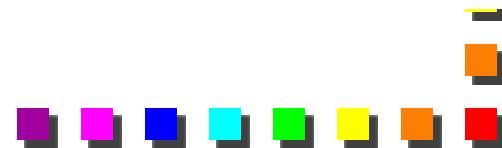
# Access Control

- ◆ Services allow setting up Access Control Lists on containers and objects
  - ◆ Accounts have one or more Identities
  - ◆ Security token encodes account+identity
  - ◆ Publically readable objects are supported
- ◆ E.g., you can set up an image or content repository for your public web site, but restrict update/delete operations
  - ◆ Generalization of WebDav



# OpenStack REST API for Storage

- GET v1/account HTTP/1.1
  - ◆ login to your account
- HEAD v1/account HTTP/1.1
  - ◆ List account metadata
- PUT v1/account/container HTTP/1.1
  - ◆ Create container
- PUT v1/account/container/object HTTP/1.1
  - ◆ Create object
- GET v1/account/container/object HTTP/1.1
  - ◆ Read object
- HEAD v1/account/container/object HTTP/1.1
  - ◆ Read object metadata



# Creating an Object

PUT /v1/<account>/<container>/<object> HTTP/1.1

Host: storage.swiftdrive.com

X-Auth-Token: eaaafd18-0fed-4b3a-81b4-663c99ec1cbb

ETag: 8a964ee2a5e88be344f36c22562a6486 MD5 checksum

Content-Length: 512000

X-Delete-At: 1339429105

Mon Jun 11 08:38:25 PDT 2012

Content-Disposition: attachment; filename=platmap.mp4

Content-Type: video/mp4

Content-Encoding: gzip

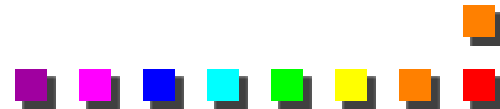
X-Object-Meta-PIN: 1234

User defined metadata

[ ...object content... ]



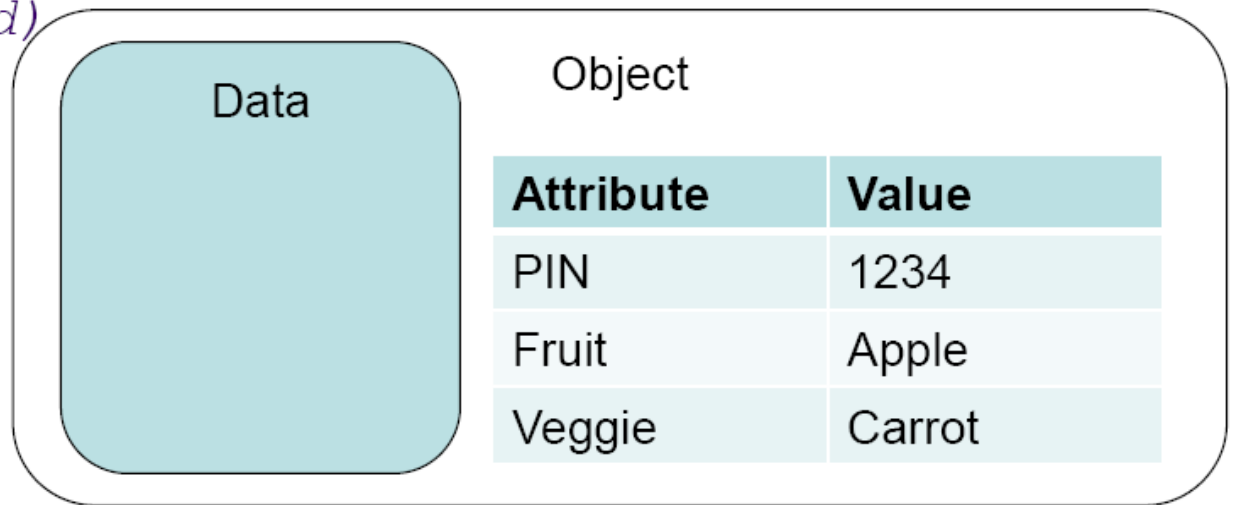
上海交通大學



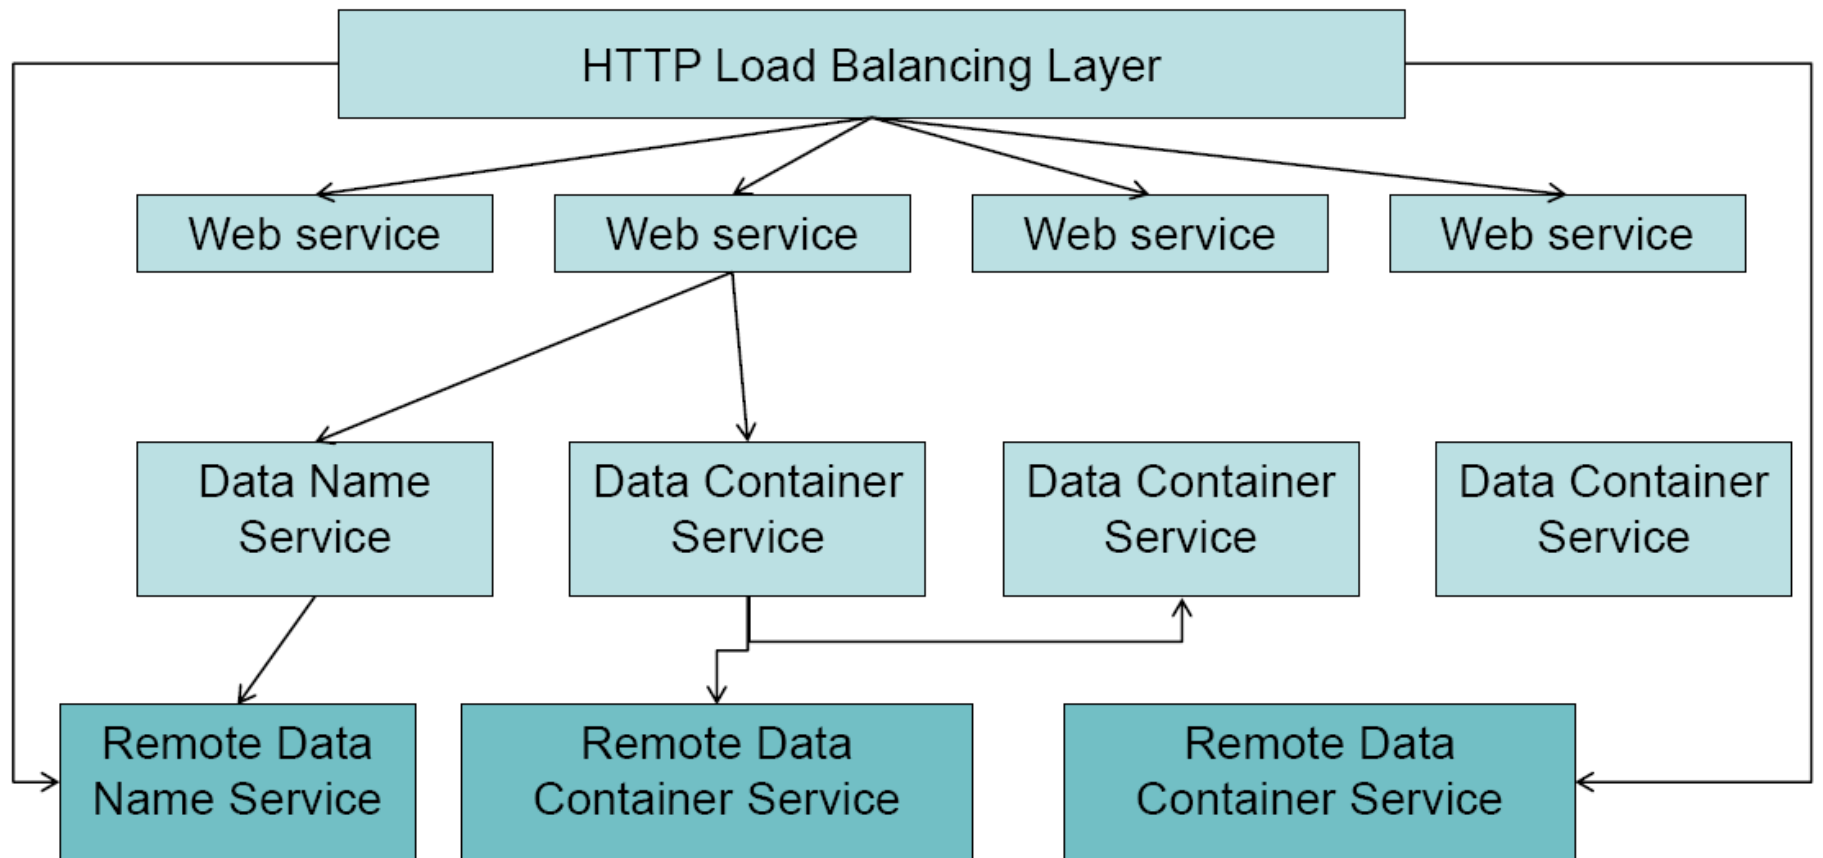
# Updating metadata

```
POST /v1/<account>/<container>/<object> HTTP/1.1
Host: storage.swiftdrive.com
X-Auth-Token: eaaafd18-0fed-4b3a-81b4-663c99ec1cbb
X-Object-Meta-Fruit: Apple
X-Object-Meta-Veggie: Carrot
```

*(no data payload)*



# Consistency and Availability





# CAP Theorem

- Impossible to provide all of these simultaneously
  - ◆ Consistency, **A**vailability, **P**artition Tolerance
- Replication across partitions increases availability
- The web access load balancer may direct requests to different partitions that may not be consistent
  - ◆ Partitions will become consistent eventually
- CAP Theorem =>  
Applications may see old versions of objects
  - ◆ Not see a newly created object
  - ◆ See a recently deleted object
  - ◆ Read the previous value of an object

# Beyond Blobs

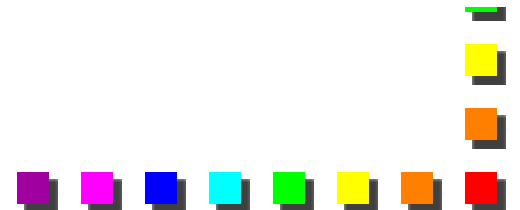


## ◆ Blob

- ◆ Write or read the whole object (PUT / GET)
- ◆ No in-place modification
- ◆ Only metadata can be modified incrementally (POST)

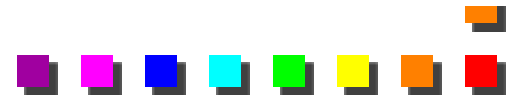
## ◆ Great for images, static page content

- ◆ Not so elegant for search
- ◆ No well defined sub-structure



# Tables

- ▶ Table has overall metadata and access control
  - ◆ Similar or same as blob
- ▶ NoSQL model
  - ◆ Key selects a row
    - › Sometimes a “partition” results in a two-part key, and makes load balancing across table servers explicit
  - ◆ Rows can have dynamically changing columns
    - › Add columns later as you decide what you need
    - › Not every row may store a particular column
  - ◆ Element at row,column has a data type
- ▶ Queries search over rows within a single table
  - ◆ E.g., find rows where a column matches a filter
  - ◆ No “join” that combines different tables



# Queue

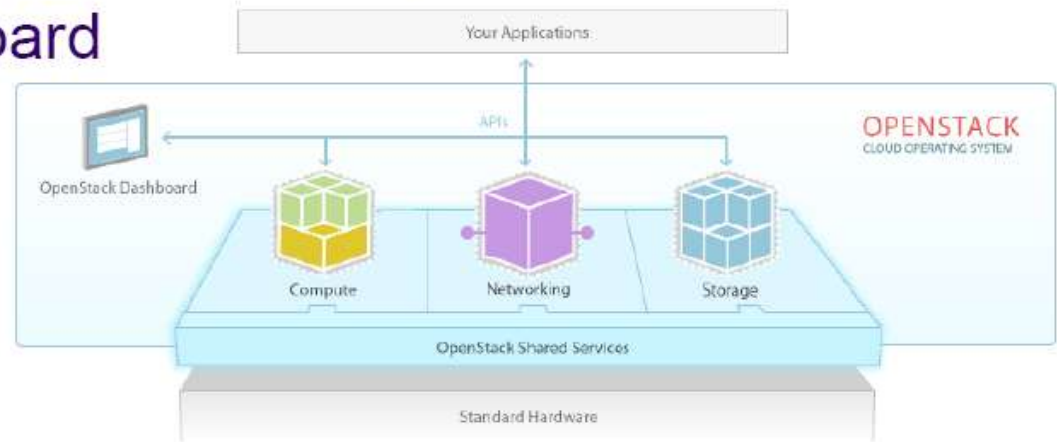
- ▶ Queue has metadata and access control
  - ◆ Similar or same as Blob
- ▶ PUT a message visible to other clients
- ▶ GET (and consume) a message
  - ◆ FIFO (First in, First Out)
- ▶ Messages limited in size (e.g., 8K or 64K)
  - ◆ Can store a reference to a larger Blob in the message
- ▶ Can Peek, Delete, and Clear messages

# Web Object Platforms

- Here we give no endorsement implied or explicit about the relative merits of these cloud platforms
  - ◆ Amazon Web Services
  - ◆ Windows Azure
  - ◆ Apache OpenStack
  - ◆ SNIA CDMI (Cloud Data Management Interface)
- These are all similar with REST APIs
  - ◆ But plenty of differences among the offerings
- CDMI is a standard for storage and storage management REST API

# OpenStack Cloud Operating System

- Open Source project part of the Apache project
- Compute (VM provisioning)
- Object Storage (in the cloud, or locally)
- Block Storage (volume provisioning)
- Networking (Software defined networks)
- Authentication (unified, multi-tenant)
- Management Dashboard

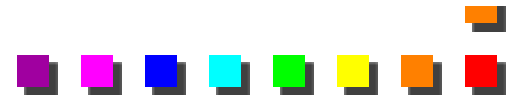




# Amazon Web Services

---

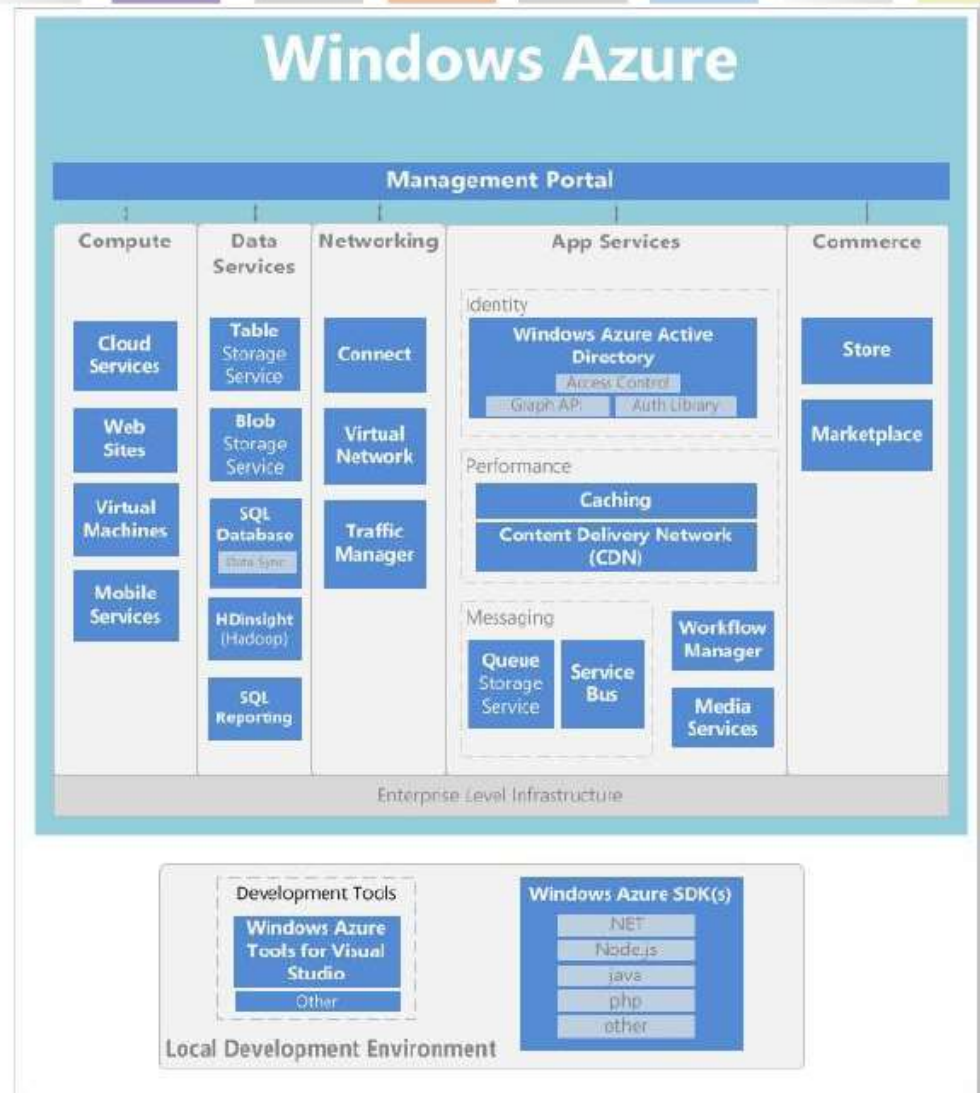
- Compute (EC2)
- Content Delivery
- Database (Dynamo)
- Deployment
- Identity
- Queue Service
- Search
- Notifications
- Monitoring
- Network load balancing
- Billing
- Simple Storage Service (S3)
- Elastic Block Storage
- Jobs





# Microsoft Windows Azure

- Compute
- Data
  - ◆ blob, table, queue
- Networking
- Applications
- Commerce
- Dev Tools



上海交通大学



# CDMI

- ISO Standard in Oct 2012
- Cloud Data Management Interface
  - ◆ Client API for blob and queue storage abstractions
  - ◆ Management APIs for accounts and access control
  - ◆ Capability APIs to introspect on system capabilities
- REST APIs for these abstractions
  - ◆ Data
  - ◆ Container
  - ◆ Domain
  - ◆ Queue
  - ◆ Capability

# Web Object APIs



## ➤ OpenStack

- ◆ <http://api.openstack.org/api-ref.html>

## ➤ Amazon S3

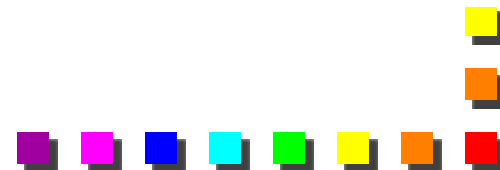
- ◆ <http://aws.amazon.com/documentation/s3/>

## ➤ Microsoft Windows Azure

- ◆ <http://msdn.microsoft.com/en-us/library/windowsazure/>

## ➤ SNIA CDMI

- ◆ <http://snia.org/sites/default/files/CDMI%20v1.0.2.pdf>



# Web Object Summary

---

- There are several competing Cloud Operating Systems that each offer similar storage facilities
  - ◆ As well as several offerings (e.g., Cloud Foundry) that seek to provide a common API
- The storage is provided in the context of a broader set of services for virtual machine management and application development
  - ◆ This is an interesting counter point to POSIX, but it is still a world of de facto standards

# Two Classes of Object Storage

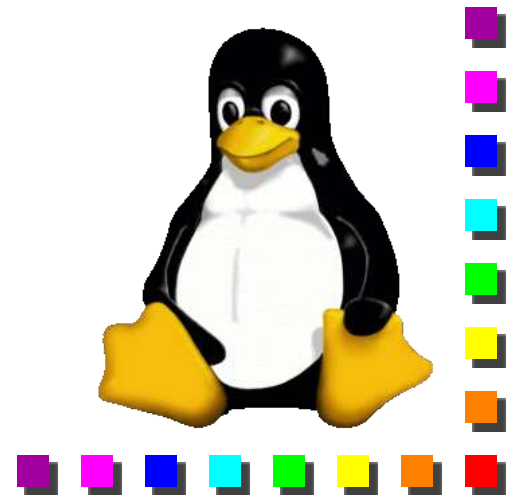
## ➤ Storage Devices

- ◆ Move smarts into the device (NASD)
- ◆ Raise the level of abstraction
  - › Containers
  - › Attributes
  - › Security
- ◆ SCSI Model
  - › OSD command set

## ➤ Web Services

- ◆ Add storage abstraction to a web-based system
  - › Containers
  - › Metadata
  - › Security
- ◆ REST Model
  - › HTTP protocol

# Project 4: File System



上海交通大學

# Source

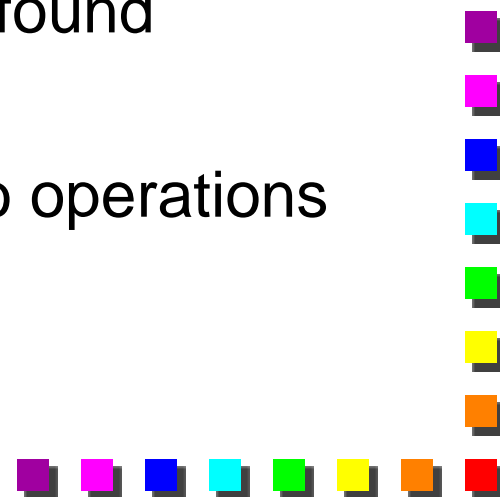
- Inode.c/Makefile (kernel source of romfs)
- Test.img (a romfs image, you can mount it to a dir with 'mount -o loop test.img xxx')
- Say test.img is mounted in t, 'find t' output
  - aa
  - bb
  - ft
  - fo
  - fo/aa





# Practice 1

- Change romfs code to hide a file/dir with special name
- Test & result
  - insmod romfs hided\_file\_name="aa"
  - Mount -o loop test.img t
  - then ls t, ls t/fo, no "aa" and "fo/aa". found
  - ls t/aa, or ls fo/aa, no found
  - Without the code change, above two operations can find file 'aa'



# Practice 2

- change the code of romfs to correctly read info of an 'encrypted' romfs
- Test & result
  - insmod romfs hided\_file\_name="bb"
  - Mount -o loop test.img t
  - Say bb's original content is 'bbbbbbbb'
  - With the change, cat t/bb output 'ccccccccc'



# Practice 3

- change the code of romfs to add 'x' (execution) bit for a specific file
- Test & result
  - insmod romfs hided\_file\_name="bb"
  - Mount -o loop test.img t
  - Without code changes 'ls -l t', output is '-rw-r--r--'
  - With the change, output is '-rwxr-xr-x'

